CHAPTER 6

DON'T RAISE MONEY FOR THE WRONG REASONS OR FROM THE WRONG PEOPLE

The best entrepreneurs . . . know how to tell an amazing story that will convince talent and investors to join in on the journey.

—Alejandro Cremades, 2016

If you want a better-than-average shot at being counted among the world's most admired, most successful entrepreneurs, you'll almost certainly need to raise money from investors (unless you're independently wealthy, or your father's last name rhymes with "rump"). Fundraising is glamorous. The press writes about you. Your friends congratulate you. Your competitors fear you. The office and the perks improve. Salaries go up. And, supposedly, everything else about your business (customer growth, recruiting, marketing) benefits, too. Plus, it can win you new friends. My wife once physically stopped me from shouting "Drinks are on me!" at a bar on St. Patrick's Day, because I'd just received a funding offer over email. Keep in mind, the money wasn't in the bank, but I was already drunk on the idea of it (and two Moscow mules).

But there are a multitude of reasons why, if you're not building a business 100 percent aligned with their model, raising money from institutional investors is a terrible, no good, very bad idea.

We raised our first round of $1.1 million in 2007 with Michelle Goldberg from Ignition Partners and Kelly Smith from Curious Office. Since then, Moz has raised two subsequent rounds totaling another $28 million. I was privileged to spend a few years on the board of another venture-backed startup, San Francisco's Minted. I was also, unfortunately, the beneficiary of months of fruitless fundraising attempts nearly every six months from 2009 to 2012. And I've spent countless hours over the years with dozens of other venture-backed CEOs and founders, hearing their stories. My exposure to the startup investment world hasn't been as deep as some, but it's been enough.

Venture capital changed Moz. It changed me. It made me a better, more focused, more ambitious entrepreneur. I learned more in my first two years as CEO of a VC-backed startup than I did in the seven prior years at a struggling family-run business with my mom.

The money we raised at Moz and the help of the partners who joined our board were remarkable gifts. I couldn't have come this far in my professional career or written this book without them. But when asked if I'd raise money again in any of my hypothetical future business endeavors, my answer has changed over the last few years, from "Yes, definitely" to "Oof . . . I really hope not."

Why? Two reasons. The odds. And the cost.

I Mean, You Don't *Technically* Sign the Deal in Blood

Founders usually think (I certainly did) that when investors put money into your company, there's alignment in the outcome. You're on the

9780735213326_Lost_TX.indd 90 7/19/18 11:16 AM

same team. Everyone's cheering for the company's success, and everyone's willing to put in hard work to make that happen.

Forgive me, but I have to take a little reality hammer and smash your founder delusion.

Look, investors aren't jerks (I mean, some of them are. Looking at you, Caldbeck. You too, KPCB.). They aren't lying when they loudly proclaim that they're 100 percent behind you and want to do everything possible to help. It's true. In the beginning. But over time, their incentives change according to your performance versus the rest of their investment portfolio. That's when misalignment occurs, and if you and your company, especially your leadership team, aren't prepared for it, reality can hit hard.

Institutional investors and angel investors alike face long odds with any individual company. That's why they place a lot of bets. A tiny number of companies will make them money, and the rest will lose money, break even, or deliver returns too small to "beat the market" (i.e., earn more than the 8–10 percent compound, year-over-year growth of the S&P 500). Realistically, the distribution for the average venture fund looks like this:

- Out of ten investments, five will fail.
- Another three will return an insignificant amount.
- The final two will combine to form the bulk of any gains.

Startup investing follows the Pareto principle: 20 percent of the investments return 80 percent of the fund.

If you go into the fundraising process, or come out of it, even successfully, thinking that you now have a partnership of equals, a partnership of aligned interests, a partnership where their success is tied to your own, you will have a shit experience. I don't want to spoil the end of the book when we're only in the middle, but trust me, it happens a lot (if you can't handle the suspense, flip to chapter 17).

Here's what I wish investors would tell founders before they put money into their company:

- "I invest in dozens to hundreds of startups. Eight out of ten don't return any money, but I don't know which ones those will be, so I have to place a lot of bets."
- "If you end up looking like one of the companies that will be that big moneymaker, I'll lavish you with attention, as will the rest of my partners. We'll make you feel important, powerful, respected—like a dear friend and close confidant, and maybe the kid I never had."
- "If things go the other way, and you look like one of the duds, expect that our attention and interest will fade; it may start to feel like meetings with you and requests from you are more of a chore than a shared mission."
- "One of our biggest tools in either preserving a growing company's prospects for success or attempting to recover a flailing startup is to replace the CEO. If things are going well, that's very unlikely. If things go poorly, especially for an extended stretch, it's much more likely."
- "If you would be happiest building a strong, stable business that's profitable, that makes you wealthy and happy, that has reasonable harmony between your work and the rest of your life, we are absolutely the wrong choice."
- "If you cannot imagine doing anything but grinding as hard as you can, with relentless focus, and demanding the same from the team around you in pursuit of becoming an incredibly rare moonshot of a billion-plus-dollar business, even though the odds suck, congratulations, our model is a match."

- "Personal happiness and successfully raising venture capital are rarely correlated."

I don't mean to be overly negative on the venture model. If I seem harsh, it's because the tech media and startup culture has done such a miserable job portraying reality. And VCs themselves—whether on Twitter, onstage at tech events, or in the coffee shops where entrepreneurs meet them—don't, either. Most will downplay the odds of failure, oversell the value they bring, and, if they're truly excited about your business, make you believe you're a world-conquering genius. For a very brief window of time, they are the ultimate optimists. And for a brief window of time, you'll be one, too.

How VCs Lose, Even When They Win

According to statistics from the National Venture Capital Association, an estimated 30–40 percent of high-potential US startups fail completely in the sense that investors lose all their money. But if you define startup success as delivering an expected return on investment, a whopping 95 percent of startups would be considered failures. That often also means they fail to deliver meaningful compensation to founders or employees, as well. Considering the countless hours, stress, and effort it takes to start a company, receive investor backing, and put that money to work, that failure rate seems almost unbelievable. It's like agreeing to exclusively produce movies starring Nicolas Cage. Who would willingly sign up for such a painful journey with such awful odds of success? And why would so many billions of dollars go into it?

(Also, have you seen *Ghost Rider*? Don't.)

To understand this quandary requires understanding the world of venture itself—how VCs get their investment dollars, what their goals are, and how their funds operate. Just as it pays to understand the motivations and attributes of a supplier, a partner, a contractor, or a potential customer for your business, so, too, it pays to know how institutional investors operate. In fact, it's actually more critical than many of those others because once you sign up for venture capital, you're usually locked in for the life of the company (or until/unless you become part of the 5 percent that meets your investors' return expectations).

Venture firms generally get money from four places: large endowments from very wealthy individuals or families; pension funds from big employers; well-endowed university funds; and evil billionaire masterminds.

(Technically, the billionaires don't have to be evil, but then they just fit into the first category.)

Together, these sources of capital are called limited partners, or "LPs." Much as startups raise money from VCs, VCs raise money from LPs. The venture capitalists will try to convince an LP's investment committee that their approach to startup investment opportunities, their team's judgment and experience, and their fund's assistance will produce a winning ROI.

The best VCs are good people as well as good investors. But they're only human and beholden to commitments they've made to their own investors. Incentives matter, and until you can empathize with theirs, it will be a struggle to understand their (mostly logical) behavior.

Their goal is to improve upon the rate of return that would have been achieved through putting the money into public stocks, bonds, or other investment vehicles. The target is 12 percent annual growth, which, over the life of a ten-year fund, means returning three times the fund size (e.g., $300 million on a $100 million fund). Beating the market is hard. Like, really, really hard. Only about 5 percent of venture investment firms actually succeed at it. A further 10 percent will return two to three times, the next 35 percent will return one to two times, and the bottom 50 percent return less than their initial fund. Thankfully, for all those employed at the 95 percent of firms who don't deliver on their goal, it can take fifteen-plus years to determine whether a fund is successful due to the extended time it takes for them to make investments and for those companies to die, sell, or IPO.

The only way the 5 percent of successful VCs make their returns is through enormous outcomes from a tiny number of companies. It's usually one or two investments out of dozens to hundreds a fund might make that deliver almost all the gains. This makes intuitive sense if you're a close observer of the startup world—for every Google, Facebook, Snapchat, or eBay, there are thousands of startups most of us have never heard of, and never will.

Moz and its investors can serve as a good example of why only "enormous outcomes" fit the model.

In 2004, Ignition Partners raised $300 million from their LPs. Those LPs expect that, over the next decade, Ignition will return at least three times the amount invested. That's $900 million.

In November 2007, Ignition used $1 million of that $300 million fund to invest in Moz. That deal valued Moz at $7.1 million, making Ignition's ownership about 14 percent of the company's shares. Let's imagine that in 2011, Moz sold to an acquirer for $40 million (a not-so-far-fetched scenario that we'll talk more about in chapter 9 while I cringe at my past decisions). The returns would look something like this:

PARTY	AMOUNT INVESTED	% OF COMPANY OWNED	AMOUNT MADE IN A $40 MILLION ACQUISITION
Rand	N/A	32.3%	$12,920,000
Gillian	N/A	32.3%	$12,920,000
Moz employees	N/A	20%	$8,000,000
Ignition Partners	$1,000,000	14%	$5,600,000
Curious Office	$100,000	1.4%	$560,000

Technically, in the scenario above, Ignition's partners are beating their model's forecast. Their investment is returning 5.6 times the capital back to the fund and doing so earlier than the LPs need or expect. But for Ignition, this would be a useless return. It sounds absurd, but it's true: following this example, the partners would need to fund three hundred companies with a guarantee of the same successful outcome on each to hit their goal.

There's no way the fund can make even half that many investments, and the success rate, as we discussed, has a nonlinear distribution that's strongly weighted to the top few companies. In the dispassionate eyes of a mathematical, dollar-focused analysis, it would be better to have put that $1 million and, more important, the time one of Ignition's twelve partners spent (sourcing the deal, making the investment, sitting on the board, and helping the company) into another startup.

Read on and tell me this doesn't sound like an insane game of high-stakes craps in a Bond film, only with nerds.

We're starting a venture capital partnership together. We'll call our firm "Scorpio Ventures"; it has a nice, subtly evil ring to it. We pitch some LPs and convince them we're investing geniuses, raising our first fund of $400 million.

After a few months of driving around in our Teslas, meeting entrepreneurs at hipster espresso shops, and crafting our Internet of Things market thesis, we find a company we love: the Globex Corporation. We negotiate with the founders, agree on a pre-money valuation of $45 million, and invest $15 million from our fund in their Series A. For the next two years, we watch them grow like a weed, supporting the team by leveraging our networks, our wisdom, and our time. At our ninth board meeting, they announce they have an offer from industry giant C. M. Burns Inc. for a whopping $450 million. The founders are ecstatic! Their company is worth ten times what it was just two short years ago. But the Scorpio Ventures team is having a bad day. Why?

Because $450 million returns only $112 million to us. That's almost ten times our original investment, which sounds phenomenal, but to reach our promise of three times the returns on our $400 million fund, we need to get to $1.2 billion. The math of startups is now working against us, because Globex was our hottest investment, and with them out of our portfolio, the odds of making up the other 90 percent–plus of our needed returns are considerably worse.

We might even be tempted to fight the founders on the company's sale, arguing that in another four years, Globex could be worth five to ten times as much as it is today. The founders' families and friends (and those of their stock-option-holding employees) might say they're foolish to give up a surefire, massive payday on the outside chance of an even bigger one years from now. After all, once you've made $100 million from selling your company, does a few hundred million more really mean all that much?

But for us, letting the sale go through (and it's up to us, because our investment terms with the founders give us veto rights over any transaction) is almost equally foolish. Upon hearing that we okayed the purchase, our own family and friends, at least the ones who understand the mathematics of venture, would be fully justified in telling us "You're totally mad, Scorpio!"

(Did I squander my youth watching *The Simpsons*? No. It was clearly time well spent.)

From this overly simplified model, anyone considering entry to the world of venture-backed startups should have at least this one crucial takeaway: unless your business is in alignment with the venture model of investing in many failures to find a small handful of absurdly successful mega-winners, this path is not for you. Get comfortable with the odds, or don't roll the dice. The venture business is about outliers.

If you want to raise money, or if you're joining a startup that's planning to raise money (or already has), you *have* to understand these odds and this risk model or you'll be an unwitting pawn in a game where the deck's stacked against you. If your startup raises VC, but you're unwilling to do high-risk things that could kill your business nine out of ten times but might make you a unicorn, you're not aligned with the venture model. That misalignment can mean losing your job and your company, or getting in a nasty stalemate with your board of directors.

Alignment means that you and your company are in this for the long haul, and recognize that only an exceptionally unlikely, multi-hundred-million- or billion-dollar outcome constitutes "success." It means you have a plan for how to turn your idea into a company with either hundreds of millions of dollars in revenue or tens to hundreds of millions of users. It means saying no to the early offers (that may or may not materialize) that could make you, your cofounders, and your early employees into millionaires so you can play the long odds of the future.

This journey, by the way, is getting longer. National Venture Capital Association data showed that the average time from funding to exit (via an acquisition or an IPO) increased dramatically from 3.1 years in 2001 to 6.8 years by 2014. Those figures count all exits, including many that did not generate a positive return for investors or founders. When EquityZen limited its analysis to only those startups that had an IPO (and thus, almost definitely generated returns for at least its investors and probably most founders, too), it found that from founding to public offering takes an average of eleven years.

In the startup world, even if you raise money, becoming an overnight success doesn't happen overnight.*

* While this chapter has focused on venture capital, there are some alternatives for aspiring startup builders. I've covered these a bit more (along with some resources) in the book's afterword.

CHAPTER 7

SO YOU'VE DECIDED TO ASK COMPLETE STRANGERS FOR MILLIONS OF DOLLARS

For the first seven years of Moz's life as a venture-backed startup, we had it pretty sweet. With growth rates of 100 percent year-over-year, operating with a model that investors and acquirers valued at many multiples of revenue, we got a lot of attention. My inbox was regularly filled with invitations to connect with people and participate in events that seem outlandish even to this day. A dinner at Sheryl Sandberg's house (thanks to an invitation from her immensely kind late husband, Dave Goldberg). An evening soiree with UN Secretary General Ban Ki-moon (at the home of one of the partners from our Seattle investors at Ignition). I'd travel to Silicon Valley and be swarmed with offers of coffees, dinners, drinks, come-by-the-office-and-let-me-show-you-arounds from companies and people I idolized. I felt important. And wanted.

But years later, when our growth rate slowed to 20 percent, then 10 percent, and it had been a long time since our big fundraising rounds, that professional and social validation from my fellow VC-backed startups waned. Our board meetings went from the most important thing on our investors' calendars, never to be missed, to getting rescheduled three times and being attended via video calls rather than in-person. The technology press stopped writing about us. Potential investors stopped emailing. Sure, our products were getting better, our company was still growing, and we were technically more profitable, but startup culture is about one thing: growth. As fast as you can. At all costs.

If I haven't scared you away from venture capital yet, don't worry—it gets worse. Because we're about to dive into the heartrending, ugly, unfair, why-does-everyone-get-it-but-you process of fundraising.

I kid! I kid! Sort of.

For some entrepreneurs, fundraising is part of the fun of building a company. They relish the exposure to critical feedback, the social aspects of those hard-to-get introductions, the prestige of being in meetings with seven- or eight-figure outcomes, and the congratulatory press and celebratory experience of finally closing the deal.

That's not me. I loathe the concept of direct sales—of convincing someone to back you or buy from you, rather than earning their interest and attention through your good works and organic marketing. Little wonder I love SEO, eh?

But regardless of your predilection or aversion to fundraising, the process is incredibly tough, with dismal odds. My goal is to improve yours by making my experiences transparent and providing a road map for how this insular, seemingly inscrutable world works.

In 2009, Moz was growing fast, profitable, and, at least in my opinion at the time, in need of another round of funding to help us accelerate that growth and dominate the SEO software market. From April to October, I talked to dozens of venture capital firms, most of them on or near Sand Hill Road, an arterial connector between Highway 101 and I-280 in Menlo Park, California, famous for its density of low-slung office buildings housing many of the world's best-known VCs. I'd spent the better part of the spring earning introductions to these potential investors, almost entirely through other CEOs who were familiar with or fans of Moz (over the years, I'd helped out a lot of these folks with SEO issues or recommendations, and they were, in turn, very kind to me).

I was twenty-nine years old at the time, and had as much impostor syndrome as I would have had pretending to belong in the Batcave. In my first few meetings I was so jittery and nervous that I can't even remember what happened (which is probably a blessing; I didn't hear back from most of those investors). But by meeting four or five, I was getting the hang of things.

You start with an email introduction, best delivered directly by one of the VC's current or prior portfolio CEOs. The one below (as near a prototype as you'll find) came via Mike Cassidy, whose company Ruba was later acquired by Google, and connected me with James Slavet, a partner at Greylock Capital, whom I'd meet that summer.

Subject: intro SEOmoz

Mike Cassidy < > 6/12/09

To James, Rand

Hey James,

I had dinner with Rand Fishkin, the CEO of SEOmoz, a few nights ago. Great guy with a very interesting business. He's been approached by a few VC's, but he's really interested in the right chemistry. I thought of you. I'll let you two chat if there's mutual interest.

Hope all is well with you!

Mike :)

Mike Cassidy
CEO, Ruba

The email connection usually kicked off an introductory phone call between me and the investor. If that call went well, they'd ask when I would next be in "the Valley," a.k.a. Silicon Valley. My 2009 fundraising attempt dictated no fewer than four separate trips to that valley, most with four or more separate meetings spread over a few days.

I'm most comfortable eating homemade cookies and watching the Die Hard canon while my wife shouts advice to the characters on-screen. This was not my world.

Just driving into the parking lot was humbling. I'd be in a bare-bones rental car, surrounded by shiny, six-figure automobiles. The lobbies were always an indescribable combination of low-key and utterly terrifying. The furniture didn't look ostentatious, but it clearly cost a fortune. The names of companies and founders on the walls weren't flashy but were meant to impress upon visitors the billions of dollars of wealth that had come through these premises. To me, these complexes were always asking "who the @#$% do you think you are, coming here?" I can't imagine how doubly intimidating it would be as a woman founder or a founder of color to see only white (and a few Asian) men on the staff, in the magazines of the coffee table, and celebrated on those walls highlighting great exits.

Usually, the partner I was meeting was running ten to twenty minutes late, so I'd wait in the lobby, willing my resolve to stay cool, calm, and confident, double-checking my notes about their bio, past companies, current deals, and anything they'd written or published. The meetings were always an hour, minus whatever time had been lost up front. If we started twenty minutes late, we went forty minutes. Maybe that alone was a bad sign for me. . . . I'd assume the more interesting deals got the overage.

But the fascinating part came at the end of the meeting or in the email follow-up from the investors who seemed genuinely interested. They'd tell me (or I'd ask) to speak to some of their portfolio CEOs and ask about their value as a partner and their support as a firm. Every VC believed himself or (in, unfortunately, only two cases out of almost fifty) *her*self to have strong references from their CEOs. I diligently followed up every time, and talked to at least two or three CEOs funded by a dozen or so venture firms. Of those, the majority did have good, though not always glowingly positive, things to say. And even though it was my first time speaking to these CEOs, many were happy to give me an honest and surprisingly cold or even directly negative feedback about their investors. I found this disconnect between the confidence of the partner and the honesty of the founder/CEO valuable and refreshing.

If you go down the venture funding path, make sure you do the same. And for those of you who get funded or who already are, please carry on this tradition of putting your fellow entrepreneurs first. It helps make the land-mine-laden game of startup investment a little less dangerous and a little more camaraderie-filled.

It's not just this closeness and rapport among founders or the reputation VCs inherit that makes them, more often than not, fair players in the startup game. Deal flow is the other component.

In the movie *The Social Network*, Mark Zuckerberg, portrayed by Jesse Eisenberg, insults a group of investors who hope to put money into Facebook, citing his friend Sean Parker, and walks out. It is indeed true that Facebook's founder arrived late to a meeting with the VC firm Sequoia (whom I pitched on behalf of Moz, twice, and whom, I'd agree, has a fairly deserved reputation for humorlessness). In his pajamas, the Facebook founder delivered a "top 10 reasons not to invest in us" pitch, and mentioned Sean Parker, whom Sequoia had fired from the board of another company (Plaxo).

That's an extreme example of a real phenomenon. VCs develop connections and reputations based on their actions, and when a founder or executive feels mistreated by their investors, word spreads fast. That mistreated founder, and their friends, are less likely to refer other startups to that investor in the future, and if a very hot deal that many VCs are interested in arises, those founders may opt to choose based on reputation and on references from their network. Thus, investors have a strong incentive, one tied to their future success in earning important and valuable deals, to treat founders and portfolio companies well.

My experience with investor Brad Feld of Foundry Group could not have been more atypical. In 2012, I had one phone call with Brad on a Monday, during which he invited me (and my COO) to come visit Boulder at the end of that same week, to pitch their entire partnership. We spent nearly the whole day with Foundry's four partners. Brad took us out to dinner. As soon as we ordered, he told us he wanted to invest. Then he spent another ninety minutes with us after the meal, talking about all the aspects of Foundry's involvement with Moz. I never felt that Brad was playing games with us, or that he ever treated us with anything but the utmost respect and friendship. He was a mensch that day, that evening, and ever since.

One conversation in particular stuck in my head. I asked Brad how Foundry would react if we were offered a deal in the next couple of years that looked great to Moz's founders and employees, but didn't provide the kind of return on investment that Foundry sought. My recollection won't be perfect, but to the best of my memory, Brad said:

> I've worked with a lot of entrepreneurs over the last fifteen years. Sometimes their first company goes well. Sometimes it doesn't. But I'm patient. I pick founders because they're people I want to work with, and I hope I get to work with them again. So if they have a chance to do what they want, we never stand in the way. My relationship is one that spans decades and companies, bankruptcies and exits. It's like a marriage—I want to be there for richer or poorer.

I've never felt so confident saying yes to an offer as I did to his.

No Such Thing as a Free Round of Funding

A question I hear often from aspiring entrepreneurs is: "How much about my company will VCs really control?" The answer is complex.

Investor and founder relations are generally positive. There are always contentious issues (as there should be in a high-growth company), but when things are going well for the company, they're also going well between you and your investors. However, if things take a turn for the worse, and it appears the company is in serious trouble, one of the primary powers a board of directors holds is to replace the CEO. Sometimes, this can only be done with the CEO/founder's consent, but other times, preferred stockholders (a.k.a. investors) have the ability to remove you or your cofounder from the company even against your will. What seemed like a minor point in a giant contract signed in a funding event can have a profound impact on your relationship to your board and your company.

I've spoken to dozens of founders over the years whose companies were "taken away from them" by their investors and boards. In some of these cases, it was probably a wise decision for the company. In others, it may have been a bad faith move that hurt the company's prospects even more. But as a founder, it's critical to keep in mind your motivations and how they align with those of your investors. Even when you all want only what's best for the company's survival, growth, and success, there can be varied beliefs about what will lead to that best outcome. Statistics are on your side—founder-led startups tend to dramatically outperform non-founder-led startups. But being removed from a leadership role, or from the company entirely, is a real possibility founders must face.

When you sign up with an investor, you'll agree to a lot of terms and clauses that can be a peculiar combination of seemingly unimportant and concerningly scary. I have three strong pieces of advice. First, do not attempt to raise money or talk to potential investors without reading Brad Feld's superb, transparent, and surprisingly fun-to-read guide *Venture Deals*. Second, get your investors to explain each piece of the term sheet (the initial offer upon which the final paperwork will be based) to you, then get your attorney (yes, you need one; no, your cousin who watches lots of *Law & Order* doesn't count) and (if you have one) a savvy entrepreneur friend to explain the same pieces to you. If the stories don't align, you'll have a strong answer to my third and final suggestion on this topic: don't sign anything with anyone you don't trust 100 percent and don't believe has your best interests (and not just their portfolio's returns) at heart.

This last one sounds more difficult than it is. Thanks to the interconnectedness of the entrepreneurial world and the general bias that startup folks have to put the interests of their fellow founders (even ones they've never met before) ahead of investors, reputation is earned and transmitted quickly. If your company fits the mold for venture, or even for angels, outreach to other founders in your field or your geography is surprisingly effective. I've been consistently amazed at how willing other entrepreneurs have been to take my emails and phone calls, to make introductions, and to give their unvarnished opinions. Take advantage of this camaraderie, and if you manage to get funding or become successful, pay it forward and help the next generation.

One final, unorthodox tip: if possible, build your expertise before you build your network, and build your network before you build your company. Each one leads elegantly into the next. If you have deep experience and skills in a particular aspect of startup building or technology that makes an hour on the phone with you deeply valuable and valued by entrepreneurs and startup teams, you've got a clear, compelling path to build a powerful network. Assist a handful of people and companies with their issues (as a consultant, a member of the team, or simply an outsider who loves to help others) and you'll have a built-in network to assist in your fundraising process. That network is what takes the fundraising process from near impossibility to potentially achievable.

Yes, it's frustrating that the worlds of startup financing are a closed ecosystem, exclusive and walled-off in their nerd-paradises. But you can make friends with the gatekeepers. And once you do, the inhabitants can be surprisingly affable.

M 9780735213326_Lost_TX.indd 109 7/19/18 11:16 AM

CHAPTER 8

FOUNDING A TOP 5 PERCENT STARTUP MAY NOT MAKE YOU RICH

Economically, you can think of a startup as a way to compress your whole working life into a few years. Instead of working at a low intensity for forty years, you work as hard as you possibly can for four.

—Paul Graham, May 2004

About once a month, I get an email from an entrepreneur telling me they're seeking investment and asking if I'd be interested in learning more about their company. My response has been the same for years:

> *Unfortunately, I can't invest in your company, because I don't have the money. Hopefully, someday, Moz's growth will provide liquidity, but for now I just own private stock. Wish you luck!*

As of this writing, my wife and I have about two year's worth of savings in the bank. She still has the used 2003 Kia Spectra she got in 2004, but I've never owned a car myself. I walk to work and back each day, often with an umbrella because we live in Seattle. Our apartment is only a few years old, in a high-density part of town, and plenty big enough for the two of us. We've even been lucky enough to help some family out with money in the past. But owning shares, even a large amount, in a private company like Moz doesn't generate the kind of payout that many folks would assume.

I don't share this story for sympathy; my salary is $220,000 a year, an amount that enables terrific freedom, the ability to help out family, and some reckless spending (mostly on travel) and helps us cover Seattle's insane rent prices. I share it because the startup culture has convinced many, many folks that if you start a company that turns into a multimillion-dollar venture, you've hit the jackpot. That conditioning has been ingrained in the gold-rush mentality of Silicon Valley geographically and of the tech startup field worldwide. But statistically speaking, this isn't the case.

Even Successful Startup Founders Don't Get Rich (Quick)

We've talked about how most startups fail. But it's not just investors who lose out. Founders and early employees are hit especially hard because they tend to take compensation below market rates for the first, riskiest years of a new venture. As is true across the economic spectrum in the first few decades of the twenty-first century, the distribution of wealth in startups goes massively, disproportionately, to the few at the very top of the field.

Conversely, a job at an emerging contender (think Box, Slack, Airbnb, Snapchat) in the tech world or at a major, public tech leader (Microsoft, Google, Amazon, Facebook) can provide salary and stock that has a vastly better chance of making you a lot of money in a few short years. Here's a comparison of the average tech salaries in Seattle from 2001 to 2015 and my own earnings at Moz during the same period:*

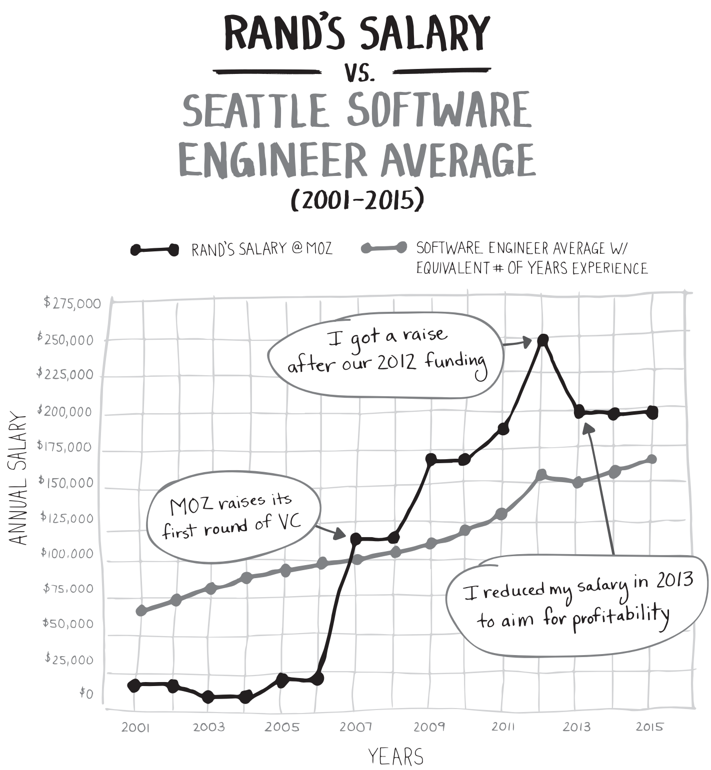

This chart doesn't account for the value of my stock, and hopefully, someday, that stock will turn into an asset that can be liquidated.

---

*In case you're curious, my salary in 2016 was about $205K and partway through 2017, I got a raise to about $220K. Seattle software engineer salaries have continued to grow, and senior engineers' total comp at late-stage startups with at least ten years' experience average just over $180K.

But many would-be startup founders and startup employees have confusion about why, sometimes, ownership of a company makes the founders wealthy versus other times when it may not.

Let's assume a scenario in which you're starting a very traditional startup—one with investors and a board of directors and the goal of a financial exit that returns money to those investors. Initially, you and your cofounders own 100 percent of the company, but once you raise money, the company will issue shares that are divided into several segments:

- Common stock: the kind you and any cofounders own
- Preferred stock: the kind your investors own (which usually grants them some special rights, like the ability to get their money out first in a sale and to have a seat on the board of directors)
- Stock options: the kind your employees own (which gives them the right to buy stock at the price it held when they were issued the option—usually when they join the company or get a promotion or bonus)

If things go well, the value of your company will increase over time, and the stock should, accordingly, have more value, too. But in reality, stock owned by you (or by your employees, should they choose to exercise their options) can only be sold when there's a willing buyer. And willing buyers are rare.

Imagine there's a supposedly valuable painting sitting in your attic. You put it up for auction, but no one bids. The painting may still have value, but it can't be turned into cash unless you find a buyer who wants it, is willing to pay, and has the liquid assets to do so. The same is true of stock in your startup. You own something that, on paper, has great worth, but it can't pay your bills or help you with the down payment on a house.

Private companies, especially startups, are very risky investments. The 90 percent failure rate keeps most types of investors away from the field. Government regulations on who can invest in startups also plays a role—in the United States, in order to put money into a startup like the one described here, you need to be an "accredited investor," meaning you have net worth of $1,000,000 or more (excluding the value of your primary residence) and/or income of greater than $200,000 ($300,000 if married) in each of the last two years.*

If the risks and the regulations don't keep buyers of your startup stock away, the fundamental attributes of private stock very likely will. When investors buy stock in publicly traded companies on an exchange like the NASDAQ or NYSE, there's a set trading price, historical data, and a wealth of information (required and regulated by law) available to everyone. In a private company, those requirements don't exist. The price can be set purely by the buyer and seller agreeing to it. No law provides buyers with a legal claim to information about the company's performance, structure, financials, or other data. Some private transactions may include information rights along with the stock sale, but these must be approved by the board of directors and add complexity and reporting requirements to the company, which can be a drag on the poor accounting/financial team (often an area with thin resources in a startup).

For a variety of reasons, investors frown on private stock sales by

*The qualifications required to be an "accredited investor," and some of the restrictions around investing more broadly are changing thanks to Title III of Obama's 2016 JOBS act. For more, see https://aaplonline.com/how-the-jobs-act-opens-deal-flow-for-non-accredited-investors.

founders (with some exceptions). Historically, they've believed that founders should be kept "hungry"; chasing that big exit (a sale or IPO) and knowing that until they do, the value of their stock is locked up. The assumption is that founders who have lots of money already will be less motivated to put all their efforts into the company's growth, and that a large payout may distract a founder as she tries to manage her newfound wealth.

Your personal success—reputation, future employability, financial liquidity—becomes inherently tied to the company's. If the company sinks, you sink. But if the company succeeds, you—depending on how much stock you own and whether or not you find a buyer for it—might succeed, too.

Sometimes, if your company is growing especially fast and multiple investors are interested, but there's not enough preferred stock in a fundraising round to go around, there may be an unusual exception made, and an investor will purchase private stock directly from founders (or, even more rarely, from early employees who've executed their stock options). But this rarity depends on a confluence of forces: an in-demand company with high growth, more appetite than availability for company stock from investors, and buy-in from the board of directors and existing investors because, even in these transactions, some special rights are often requested by the new shareholders (e.g., board observer seats, quarterly informational updates that normally go only to board members, etc.).

But if your startup is making millions of dollars and growing, could you, as founder and CEO, bump up your salary or your bonus and profit that way? In short, no, at least not without the consent of your board of directors. The board sets the CEO's salary and approves the salaries and stock compensation of employees, as well. This compensation is, in almost every startup, a function of the market averages. Worse, if you own lots of stock (even if that stock is illiquid), it's factored into your "total comp," and thus serves to make the cash portion of salary your leadership or board sets even lower.

(I have spent long hours trying to explain this to my brother-in-law to no avail. He repeatedly asks to borrow a couple million. Sorry, Ed. We really don't have it.)

Your investors have access to a diverse set of salary ranges across their portfolio companies, and they often buy additional data from aggregators of salary information, too. Using these aggregated statistics, they'll help you determine compensation ranges for the company. Reasonable arguments can be made on either side: being thrifty and relying on passion and the promise of stock option values to incentivize employees versus ponying up the top echelons of salaries to compete for talent with big companies and other high-paying startups in your field. But arguing for a salary that's far outside these ranges for a CEO or founder will garner contention and almost certain disapproval from your board.

It Takes a Lot of 7s and a Long Time to Win at Startup Roulette

In essence, founding a startup or being an early employee means taking a risk. You're sacrificing the certainty of a potentially higher-paying job with greater benefits at a more established company for what is typically a less-than-market-rate salary, bare-bones benefits, and the hope that if things go very well, and your company is one of the few that survive and thrive, your compensation will rise, the benefits will get better, and someday, if you're *extremely* lucky, your stock's value will exceed the delta between what you could have made in another job.

It's a gamble. But because the stories we hear always focus on the big winners—the Zuckerbergs and the Andreessens—we tend to forget that.

We also forget that startups take a long time to exit. Much longer than what popular culture would have you believe. Moz itself started in 2004. Thirteen years later, it is still going strong but despite moving on from what most would call "early-stage" to "mid-stage," has neither graduated out of the startup need for rapid growth nor returned money to its investors or employees. In many ways, Moz is a lucky outlier: the majority of startup founders' efforts fail entirely. But we are a great example of how foolish the idea that you can "invest four years and become a millionaire" is.

For all these reasons and more, it's often the case that founders of startups, even those that appear, from the outside, to be massively successful—raising large amounts of funding, growing at a fast clip, making a lot of money—may be (at least until or unless an exit arrives) financially worse off than their counterparts working mid-level jobs at a traditional firm. The venture capitalist and blogger Mark Suster gives wise advice to founders on how to talk about options with their team:

> *We give out stock options. I hope they're worth money to you some day. But let them be icing on the cake. If they pay off handsomely that's great. But don't count on it. Don't let it be your motivator or your driving decision.*

This same advice should be taken just as strongly by founders themselves. Your stock's "value" can blind you to the reality that private company stock is very difficult to sell, rarely appraised similarly to your investors' stakes in your company, and, if that sale ever does come to pass, it's reliant on forces largely beyond your control.

I've met plenty of people who've taken pay cuts to work at startups and then found out their stock is worth nothing. (It's happened to my friends and colleagues, and even to my wife.)

Does this mean you shouldn't found a startup? Or join an early-stage company? No. But it means you almost certainly shouldn't do it exclusively for the promise of short-term wealth. The idea that you can compress an entire career into a few years of work and be directly or consistently rewarded with cash for that compression is insanity borne out over and over by the stats. Yes, startups have better odds than the lottery, but they're dramatically worse than "put it all on red" at the casino. The myth that "founders get rich" has brought thousands of people into the world of startups potentially for the wrong reasons and almost certainly with the wrong expectations.

There are logical, wonderful reasons to start a company or to join a risky, early-stage venture. Chief among these is the freedom early-stage companies provide, especially to founders, to determine what you work on, how to structure that work, who to hire, who to keep on the team, and how every aspect of your organization will operate. Plus, you might have an idea, a product, or a mission that you really want to share with the world. It's an immense, stressful responsibility, but it's also intensely rewarding when it works.

And, many times, even when it doesn't.

One of the great things startups can do is massively accelerate your career path. If you've felt trapped in positions that don't challenge you, or constrain your earnings or influence potential, a few years at a startup, even one that ultimately fails, can dramatically shift that reality. Early-stage companies need people who are self-motivated, mission driven, and can get immense amounts of work done in short periods (by working very hard, by being efficient with their time, or both). Demonstrating that you have strategic vision, the ability to execute, and the qualities needed to recruit, motivate, and lead will make you a rare, desirable commodity in the modern economy. One of the reasons so many early-stage companies sell in small acquisitions (often called "acquihires") is precisely to get these proven, self-driven, multitalented people on board.

Moz has completed six transactions to accomplish precisely this. One of those was a young man who'd just graduated college but had built a useful product in his spare time at school that proved to me, and to our team, that he could be incredibly valuable. We paid a small acquisition price ($18,000) and brought him on board at Moz at a salary, with stock options, and with influence greater than what we'd have offered a candidate who simply applied to a job posting. Another was a pair of SEO professionals who'd built a successful consulting practice and whose skills we wanted internally for our product and engineering teams. We paid $330,000 (plus stock options and retention bonuses) to get them here, a nice multiple on their business in addition to strong salaries and benefits.

A startup can be a great way to visibly multiply what companies are willing to pay to bring you aboard. It might be a great way to level up your skills. And it has the outside chance of making you remarkably wealthy.

Just don't go in blinded by the money. Most of the time, startups are a comparatively poorly rewarded labor of passion.

CHAPTER 9

SCALABLE MARKETING FLYWHEELS > GROWTH HACKS

Growth hackers are a hybrid of marketer and coder, one who looks at the traditional question of "How do I get customers for my product?" and answers with A/B tests, landing pages, viral factor, email deliverability, and Open Graph. On top of this, they layer the discipline of direct marketing, with its emphasis on quantitative measurement, scenario modeling via spreadsheets, and a lot of database queries.

—from Andrew Chen's "Growth Hacker Is the New VP Marketing," 2012

At the start of 2009, in the wake of the great recession, Moz was in remarkably good shape. Following the successful launch of our link index tool, the team was frantically working to double down on our progress. But costs were holding us back. With greater resources, we could grow our data set and ward off any competitive pressures. Without it, we feared a well-funded copycat could easily take the market from us.

I wanted to raise another round of venture capital to help us scale.

I believed that, despite the awful VC climate (investment dollars going to startups had fallen off a cliff after the financial crisis of 2008), investors needed somewhere to put their money, and Moz was a growth opportunity with a steady two years of rising subscribers and revenue. In the first half of 2009 I made multiple trips to the Bay Area from Seattle, spent thousands on flights, hotels, and rental cars, and, after pitching more than forty individual partners at VC firms, had nothing to show for it.

So instead, we turned to the now-ubiquitous "growth hack" model in hopes of boosting revenue. That growth hack, in the form of an email marketing campaign, brought in a tremendous number of new customers and new revenue. But over time I've come to wish we'd never done it. At least, not the way we did.

It's Just a Little Hacking; What Could Go Wrong?

Our hack started with the assistance of a pair of brilliant marketers out of the UK, Ben Jesson and Dr. Karl Blanks, cofounders of Conversion Rate Experts. Their specialty was taking poorly performing web pages designed to sell a product and massively improving the percentage of visitors who converted into customers. This practice, called "conversion rate optimization," or CRO, is a powerful, meaningful part of a web marketer's toolbox. It's easy to see why: if you improve your conversion rate it has a massive impact on your customer and revenue growth (e.g., for every 100 visitors who visit your site today, 1 buys something, but tomorrow you make changes that double that to 2, or heck, just boost it to 1.1).

Ben and Karl, alongside an equally brilliant employee of theirs, Stephen Pavlovich (who'd later go on to found Conversion.com and marry one of my employees—long story, the wedding was lovely), worked with Moz to design three things: an update to our website's home page, a new version of the page that sold our software subscription, and a promotional email campaign. Their process sounds simple but is both ingenious and remarkably effective. I implore you to copy it:

> **Step 1:** Ben, Karl, and Stephen asked us for contact information of three different types of Moz users:
>
> > Paying subscribers
> >
> > Subscribers who'd tried Moz's products but left
> >
> > Members of the Moz community (who engaged in our blog comments and discussion forums) but who hadn't yet tried our subscription
>
> **Step 2:** They conducted phone interviews (and a few in-person interviews at conferences and in their offices) with several dozen members of each group, asking questions like:
>
> > What do you do professionally? What's your job title? What are your responsibilities?
> >
> > What made you initially sign up for Moz? What objections did you have, and how did you overcome them?
> >
> > (For those who were long-term customers) What do you use Moz for? What does it help you do?
> >
> > (For those who'd signed up, but quit) What did you hope Moz would do that it didn't? What made you cancel your subscription? What would have made you stay?

(And, for those who were free members of our community, but had never signed up for our software) What holds you back from signing up? What would make you try Moz's subscription?

Step 3: They used the answers to identify the right target customers for our products and to craft messaging toward this group. For us, these tended to be professional web marketers who cared deeply about their search rankings and search traffic. They were either consultants (independent or at an agency) or in-house marketers (who worked full time for a single website/brand).

Step 4: They compiled a list of objections that those who'd frequently visited Moz but never tried our software had to signing up. They additionally built a list of reasons subscribers loved the product, descriptions of how they'd overcome objections of their own, and a long set of testimonials based on those interviews.

Step 5: Stephen worked with us to design a new landing page that focused on addressing the objections we heard most commonly among the group that matched our best customers in terms of their traits and professional focus but had never signed up. That new page was almost eight times the length of the original version. You can actually still see a comparison of the two on Conversion Rate Experts' website via a case study they made.

Step 6: Lastly, we all worked together to create a promotional offer that would go out to all the members of our community who'd never tried our software before. This list numbered nearly 120,000 email addresses.

The new landing page was the first big win. When compared with the prior version, it converted visitors into buyers at nearly twice the rate of the prior page, a phenomenal improvement. To this day, I'm a huge believer in the power of Conversion Rate Experts' objection-gathering and objection-addressing methodology. It's something I urge marketers of all stripes to attempt on their own landing pages.

But good conversion practices don't fall under the "growth hacks" umbrella. Our email campaign, however, did. The original email from 2009:

Hi [redacted],

Thanks for hanging out on the SEOmoz blog this year; I'm thrilled you're a fan of our work. As a special thank you for your support, here's a gift that will (in my humble opinion) have an enormous, positive impact on your SEO performance in 2009: **a full month of SEOmoz's PRO membership for only $1**.

There's only one itsy bitsy teeny catch: because we offer one-on-one Q+A with the SEOmoz staff, we've had to limit the number of places available at the discounted rate. So while we're sending this offer out to 122,451 SEOmoz members, it's only valid for the **first 5,000 people who respond**. Don't delay—we'll be promoting our once-in-a-lifetime $1 offer on the blog on Monday, February 9th. So act now, before the riotous, can't-be-tamed masses hear of this.

To claim your first month of PRO membership for just $1, visit www.seomoz.org/trypro and enter **SUCCESS09** as your promo code. The code expires February 13th (that's next Friday), but remember, space is limited.

IF YOU DECIDE YOU DON'T WANT TO CLAIM THIS SPECIAL $1 OFFER . . . then please send a reply to this email with a brief explanation of why you aren't interested (and don't worry about hurting my feelings; my wife says it "builds character").

I hope you have a prosperous 2009!

Thanks,

Rand

P.S. Here's the link again, just in case you missed it :) To claim your full month of PRO member services for just $1, visit http://www.seomoz.org/trypro (using SUCCESS09 in the promo code!)

In classic "growth hack" format, we made an offer with limited quantity to leverage scarcity bias, used a massively discounted price that was way below our usual rate (at the time, $79/month), and time-boxed the promotion with an expiration date. If the email sounds a little like one of those "limited-time offer" TV infomercials, that's no coincidence. The same tactics apply to both.

On Wednesday, February 4, we sent out 122,451 emails with the headline: "Try SEOmoz Pro for Just a Dollar" . . . and the reply-to email address was my personal one. (The result was, among other things, a seven-hour email marathon in which Geraldine and I manually replied to more than two thousand messages. My poor wife remains uncompensated for many of the hours she poured into the company; sorry, honey!)

That was an intentional move, designed to improve the delivery rate and to make it a more authentic offer. But it also resulted in thousands of email responses over the first twenty-four hours, most of them asking exactly the same question—"Am I obligated to pay for the subscription after the first month?" As a result, we quickly decided to send a follow-up email with the title “I made a mistake about the $1 offer email,” clarifying that no lengthier obligation was required, and subscribers who paid us $1 for the first month wouldn’t have to pay anything else in the future.

That email had an even higher engagement rate than the first one, and drove a huge number of visitors and new signups. All combined, the two emails and a blog post the following week drove almost the full five thousand new subscribers we said we’d take, more than doubling Moz’s paid membership. By our estimates, the email offer and the resulting signups drove about $1 million in additional revenue for the business.

It wasn’t until later that I learned why growth hacks deserve the name “hacks.”

They’re Called “Hacks” for a Reason

The email offer didn’t make our product better; it didn’t make our subscription stickier; it didn’t help people to do their jobs better. It simply created a short-term boost of attention that led to a lot of complex, long-term problems. Among those:

- The subscribers who signed up via the $1 offer had a much lower retention rate than subscribers who’d signed up via a non-promotional offer. For years after the promotion, the remnants of those five thousand promotional signups would haunt our churn numbers (one of the most important metrics for a SaaS business).
- As a team, we became overly enamored with the impact this email campaign had on our revenue and growth. We spent years trying to replicate it through a series of aggressive discounting and limited-time-offer tactics, with middling success. In retrospect, an addiction

to finding the next great "hack" dissuaded us from the long-term product and marketing investments on which we should have focused.

- Promotional pricing, especially when offered to a very large audience, creates an impression that discounts and special offers are part of a brand's ethos, and that rather than signing up at full price, you, as a potential buyer, should wait until the next promotion launches. We found that a large swath of the SEO community viewed Moz this way after our ongoing offers (usually two to three every year). If you've ever stopped yourself from paying full price for a brand or at a store because you suspect there'll be a sale in the future, you're familiar with the mind-set.

The conversion rate optimization efforts on the landing page for our subscription product, however, were an ongoing success, and something we've continued to invest in with positive long-term returns. The CRO efforts of identifying the right audience, discovering their objections, overcoming those objections with information on your landing pages, and smoothing the checkout/signup processes aren't "hacks"; they're improvements to the fundamental flywheel powering our marketing.

I deeply empathize with the temptation to chase growth hacks. Startup marketing blog posts and presentations are filled with stories of how one great tactic transformed the growth curve of a nascent business and made them leaders in their field. If you've spent any time investigating this realm, you've probably heard the same stories.

Airbnb's hack was to scrape Craigslist's vacation and rental home listings, then contact all of the owners and convince them to also list their properties on Airbnb (or, according to some accounts, do so without even getting permission). Technically, this violated Craigslist's ToS (Terms of Service—the publisher's guidelines for how the website is allowed to be used), but it's now the stuff of legend among growth hackers of the startup world. They point to this story and argue that you can't make a growth omelet without breaking some ToS eggs.

Dropbox's hack was a double-referral system wherein, when a Dropbox user referred someone else to the service, both the referrer and the receiver of the referral got upgraded account benefits. For years, that growth "hack" was featured on marketing event stages and in blog posts across the web. The endless copycatting of this tactic resulted in plenty of frustrations as other businesses failed to capitalize the way Dropbox had. When Drew Houston, Dropbox's founder, presented his "Startup Lessons Learned" in 2010, he wisely noted that "marketing tactics for one market type fail horribly in others." Dropbox itself had tried to replicate the growth hacks of other startups without success. Drew's deck walks through these many failures in paid search, PR, affiliate marketing, and AdWords.

Hotmail's hack was one of the earliest in the web field and wasn't referred to commonly as a growth hack until many years later. If you were on the web during the late 1990s and early 2000s, you'll probably remember it. Hotmail's free email service featured a line at the bottom of every email sent proclaiming that the email had been sent through the free service and inviting anyone who received the email to sign up for their own account. The rapid expansion of the web, and of email itself, led to millions of people taking advantage of this message in an era when many other email services charged a monthly or annual fee.

In 2006, Yelp launched a hack that had an immense benefit to their overall traffic and branding—the website badge strategy. Yelp sent restaurants with four- or five-star ratings from their users a visual

9780735213326_Lost_TX.indd 129 7/19/18 11:16 AM

"badge" restaurant owners could put on their websites to showcase their positive reviews. The badge linked back to the Yelp page, sending Yelp both traffic and, through the search engines' love of links, high search rankings for restaurant names, categories, cities, and more. That strategy had been tried before, by TripAdvisor and Citysearch, but no one nailed the process more effectively than Yelp, and for years, Yelp rode the wave of SEO returns those badges helped enable.*

These stories aren't atypical. PayPal's $5 signup referral, Uber's many city-by-city hacks (the littering of referral cards, the unpaid posters in bar and restaurant bathrooms on Friday nights, the totally evil practice of faking ride pickups to their competitors so as to hurt their profits and response times), Facebook's college-focused growth tactics (e.g., get the Greek systems on Facebook and everyone else will follow) in the early days, and dozens more are pointed to by startup founders, investors, and pundits as evidence that finding the right, innovative "hack" has replaced classic marketing practices as the way new companies can and should achieve sky-high growth rates.

Like most stories of success, there's a kernel of truth surrounded by a mountain of hyperbole and oversimplification.

Some growth hacks do work. Most don't. Even at the companies just described, the vast majority of individual growth tactics didn't take. Sadly, they don't get the press coverage or focus. The lesson should be: these companies tried dozens of innovative marketing tactics, combined them with strong, constantly improving products and plenty of traditional marketing best practices, and in some cases, a

---

* If it sounds like I'm unusually familiar with this hack, that's because it was my idea. . . . Weird, right? I worked as an SEO consultant for Yelp at the time, and the badges with embedded links back to Yelp's site was a great way to improve their rankings in Google, so I suggested it; they approved and executed. Additional credit to Michelle Broderick, who worked on the project from Yelp's side.

few specific tactics proved particularly effective as part of this mix. Instead, the shiny-object-chasing narrative is the one that earns headlines and lives on in the subculture of technology startups.

But if observing the life cycle and hardships of company formation and growth has taught us anything, it's to look beyond simple explanations for the deeper, more complicated truth. We should do that here, too. Growing your startup's brand, customer reach, conversion rate, retention, engagement, and virality *can* include finding that one great hack, but to do that, you need a broader understanding of the problem you're working to solve.

The Alternative: Sustainable Marketing Flywheels

Great companies are almost universally fed by a powerful, ongoing set of marketing processes that earn attention from the right audiences and bring them to the company's (physical or virtual) doorstep.

I like to describe the complexities of the marketing process as a flywheel. This metaphor refers to a piece of machinery from the industrial revolution that stored up rotational energy from inconsistent sources in the form of inertia. This could then be used to power any number of systems that require a consistent output of that energy. I'll show you how Moz's marketing flywheel works as an example.

Our flywheel is powered by content that audiences in SEO find through a variety of channels: search engines, social media, word of mouth, conferences and events, email subscriptions, referring links on other websites, etc.

The process of creating that content, amplifying it through various channels, reaching new audiences via that amplification, and bringing people back to our website is a powerful, ongoing system that drives millions of visitors and thousands of new software free

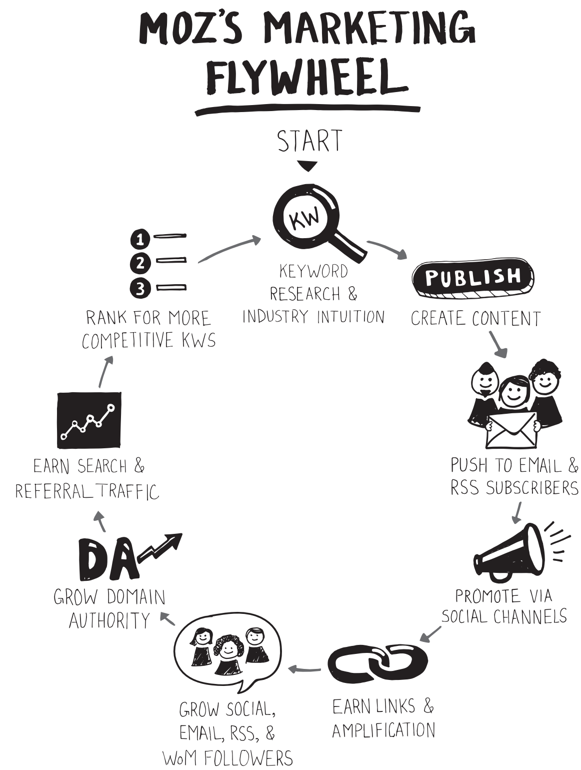

trials each month. But like a flywheel, it took an immense amount of energy to get started, and only after it was rotating smoothly, growing its inertia, did it function in this friction-light fashion.

For the first five years of Moz's existence, I blogged four or five nights a week. I'd take my intuition and my experiences from client work, online forum participation, conversations with other SEO professionals, and news coming out of the search engine world and turn it into content of all kinds. I made written blog posts, illustrated visuals, PowerPoint presentations, live videos filmed against a whiteboard, webinars, statistical surveys, interactive quizzes, and more, and I put them on our website. In those first two years, I was lucky if anything

I created received more than a few dozen visits. In 2006, after the success of the Beginner's Guide to SEO, that trajectory improved and I could regularly reach several hundred folks with new content, then a thousand, then more.

The effects built on one another. The more people were exposed to something we'd created, the better chance it had of being amplified. Those amplifications, often in the form of links and shares, led to better search engine rankings, which drove more traffic and exposure and generally more-qualified visitors (no surprise that searchers, who are actively seeking out a specific answer or resource, engage at higher rates than those who've simply seen something of potential interest in their social, news, RSS, or email streams). This flywheel has powered our marketing efforts for years. It brings to our website high quantities of visitors interested in search engine optimization and web-marketing topics, and through our content, we hope to instill knowledge of and trust in the Moz brand. Later, if and when those same people are seeking out software and tools to help with their SEO efforts, we hope to earn their business and often do.

In fact, an interesting stat about Moz's marketing and customers was uncovered a few years ago as we analyzed our visitor and conversion funnel. These numbers are somewhat older now, and have likely changed somewhat, but retain this characteristic:

Say you visited Moz for the very first time via a Google search for a phrase like "SEO tools," then immediately signed up for a free trial of our software. Chances are good you'd be a Moz customer for less than four months versus the overall global average of about nine months.

But if you visited Moz twelve times or more in a three-month span before signing up for that free trial, chances are you'd stick around for fourteen-plus months as a paying subscriber. Whoa. I know.

Turns out, our best, most loyal customers tend to be those who've spent considerable time on our website, participating in our community, consuming our educational resources, and testing out our free tools. Thus, it's actually in Moz's interest *not* to promote our products or conversions too heavily or too fast, especially to new visitors. The classic funnel optimization promoted by many marketers has this peculiar idea that we must race to turn as many visitors as we can into paid customers and that any missed opportunity represents a flaw in our marketing process. Our metrics show just the opposite. If we want to have the best long-term impact on customer growth and retention, we need patience. We need to wait for our audience to be ready and engaged with us before we nudge them toward our subscription.

I think this foments a beautiful, symbiotic relationship between our values, our content, and our paid subscriptions. We want to help people do better marketing. We want them to learn first and sign up only if our products are right for them. And our business actually prospers most from long-term, low-churn, high-engagement customers—the kind we get from deep investments not just in a visitor's conversion path, but in their professional and educational journey.

Pro Tip: If you have any type of subscription or recurring revenue, make sure you measure LTV (Lifetime Value—the total revenue customers spend during their relationship with your firm) by referral source(s) and by the number of visits prior to conversion. If your stats look like Moz's, you'll probably want to adopt a similar, slow-burn conversion process.

The power of the marketing flywheel is clear to us. But it's not just Moz for whom this works. Of the most successful startups, nearly everyone has a clearly identifiable marketing flywheel that brought awareness and traffic from the right audiences and helped those people convert to a sale or a signup at the right time.

Dollar Shave Club, the famous Los Angeles startup that offers traditional men's razor blades for a few dollars a month (originally one dollar, until their acquisition by Unilever), built a funnel based on humorous, online videos that positioned them against stodgy, expensive shaving product companies. Those videos would earn massive viral views from an audience perfectly poised to help them spread (usually young, heavily online, social-media-savvy men). The videos earned news coverage, which themselves got shared, and all of it combined to create a mass of traffic that Dollar Shave Club then bought remarketing and retargeting ads against (along with earning loads of high rankings for key search terms in Google).

Zillow built a remarkable flywheel on the initial strength of their home price calculator and "Zestimate" (Zillow's patented formula for estimating a home's value), and the traffic, shares, links, stories, and controversies that inevitably followed. Initially, visitors weren't going to the site to buy houses, but once they began associating Zillow with residential real estate data, and once Zillow leveraged the engagement, content, and links to help earn high rankings in the search results, the outcome was inevitable.

WP Engine, the popular WordPress hosting site renowned for its reliability, started by writing about and serving WordPress sites that had been overwhelmed by viral traffic, often to the point of going offline. By associating their brand and technology with the "Reddit Hug of Death" (and other, similar phenomena that overwhelm a site's web servers by sending huge amounts of visitors in a short window), and by appealing specifically to site owners whose domains had received this sort of popularity (as well as the tech-savvy crowds that often influenced them online), WP Engine developed a cult following that led to broader popularity and traffic over time.

While flywheels are critical to build, there is overlap between investing in them and experimenting with growth hacks. My experience has been that the best time to leverage a hack is when it perfectly fits with an area where your flywheel is experiencing friction.

Early on, Moz's flywheel struggled most with how to get in front of bloggers, journalists, website owners, and influencers who might be likely to link to our content. I recall vividly how, in 2005, when a few websites linked to a piece I'd written, I stood up at my desk cheering and bought a bottle of cheap champagne on the way home from work. It might sound silly, but in those days, I knew that links to my posts were what stood between toiling in relative anonymity and earning highly relevant visitors who might become consulting clients. The friction in our flywheel was how to get in front of those who might link to us and how to convince them to create those links.

The solution was to find a hack—the right hack—that could get us noticed by these specific sorts of website owners and bloggers and writers in the search technology and marketing fields. We found our hack with a piece of content we called "The Search Engine Ranking Factors."

To be fair, we were not the first or only website to attempt to list out the elements Google used to determine how sites and pages ranked via their famed algorithm. But we were the first to do it with a community-recruited crowd of influencers. Our approach relied on crafting a survey of all the various potential ranking inputs, then asking notable professional SEOs to take the survey and contribute their opinions. The resulting document aggregated quotes, summed up and averaged numerical rankings, and then ordered them by relative weight. More than a hundred folks with websites and followings of their own gave their input, and each received a personal email from me, thanking them for their participation and asking for their help in spreading the work.

This growth hack produced high value content—the ranking factors document itself—and a long list of people who helped us overcome our biggest challenge: earning links and amplification. Because each influencer had contributed, they were predisposed to help share the content. Our drought of awareness was over, and within months, nearly all of the websites where our respondents posted their own works had referenced ours. Not only that, but those links had the hoped-for impact. Moz (at the time using the SEOmoz.org domain) ranked number one for "SEO ranking factors," "Google ranking factors," "search engine ranking factors," and a host of related keyword phrases that collectively drove thousands of monthly searches.

We identified our flywheel. We found the point of friction. And we applied a growth hack to ease that friction and let the wheel start to spin faster. Today, this tactic of including influencers in the creation of content (often called "roundups") is a staple of the content marketing practice. I'd even say it's massively overdone at this point. But a hundred new, creative opportunities await.

When thinking about how to build a marketing process that's going to work for the long term, that can scale without friction, that can build on itself even as your business grows, consider the flywheel analogy. Each one will be different, and yours should be substantially unique from your competitors, built to take advantage of your particular skills, and targeted to your specific audience.

The "hacks" or marketing tactics you employ should be in service to this funnel, not instead of it. If you've got a great idea for a landing page or a referral program or a way to reach the right customers via a social network in a scalable manner, just make sure you know how to test, track, and apply it inside the funnel you're building. Growth hacks alone can't solve all your marketing problems, but the right ones may add immense value to an already humming marketing flywheel.

CHAPTER 10

REAL VALUES DON'T HELP YOU MAKE MONEY (IN THE SHORT TERM)

Corporate values, usually chosen by senior executives, are adopted to prevailing business circumstances and are not rooted in fundamental philosophical convictions, morality or ethics. In this sense, corporate values are often selected as a strategy to "rally the troops," and therefore, manipulative in nature.

—Ray Williams, 2010

My friend Rob Ousbey once had a brilliant idea for how to make Moz millions of extra dollars each year. So why did I tell him, "Hell no, we're not doing that?" Read on, friends.

In a SaaS business like Moz, churn is perhaps the most important, most studied number. It shows what percent of your customers are canceling their accounts each month (or year). A high churn rate means you need to earn a lot of new customers just to make up for the ones who are leaving your service. A low churn rate means that growth is vastly easier and usually much faster. Investors, in particular, tend to value SaaS businesses like ours with low churn rates at far greater multiples than those with high churn.

In 2011, our churn rate was about 8.5 percent per month and we had about 10,000 subscribers. This meant that, each month, we had to find 850 new subscribers to sign up, just to maintain our revenue. If we wanted to grow, we needed even more.

We were, in fact, signing up far more than 850 new customers each month, which enabled the business to scale quickly despite the high churn rate. But we were scared about the inevitable consequences of such a brief customer engagement. On average, folks stayed with their Moz subscriptions for about eleven months. Based on some research, we estimated there were close to a million potential Moz customers in the English-speaking world, but it's not hard to imagine how, in a few years of thousands of monthly signups, we could work our way through a significant portion of that group and put the long-term future of the business in jeopardy.

We needed to make our subscription stickier and provide value to our customers such that they'd want to stay with us for multiple years. That's where Rob came in.

"Why," he asked, "do you let people cancel with a click right on the website?"

"You think we should make it harder to cancel?" I replied.

"If you change the cancellation process to a phone call," Rob explained, "I'm willing to bet that friction alone will improve your monthly churn rate. Plus, you can then talk to folks as they're canceling and get a much better understanding of who they are and why they're leaving. You could probably save a good number of them, too, or get them back as customers again in the future. Maybe switch some to a cheaper plan or convince them to buy a different product from you."

I'd seen some stats that backed up Rob's guess and told him so.

"You're right. Making the subscription cancelable only by phone can have a real impact on churn rate. But we all hate services that require a phone call to cancel after a purely online signup. They're intentionally making it difficult. It's not empathetic, and that means it's not TAGFEE."

"Fair enough," Rob replied, "but you're leaving a lot of money on the table."

"Well," I answered, "we always say: they're not core values if you're willing to sacrifice them in exchange for money."

Yes, but Is It TAGFEE?

Values may not make you money in the short term, but they're invaluable to any business in the long term. Values are not always easy. They force hard decisions. They can work against short-term growth. They restrict paths that might otherwise be open to pursuit. Establishing and adhering to core values carries great intrinsic and extrinsic benefits. But usually these become evident over the long term, and that can be immensely frustrating for startups struggling simply to stay alive long enough to get to profitability or fundraising. This tension is hard, but my experience and the correlation of values adherence to performance suggest it's worth it. Plus, sticking to values makes it possible to look in the mirror without hating the person staring back.

At Moz, we have six core values, represented by the acronym TAGFEE—Transparency, Authenticity, Generosity, Fun, Empathy, and the Exception. These are the beliefs we prioritize above the success or growth of the business. TAGFEE acts as a litmus test for whether we should or shouldn't take an action, hire or let someone go from the team, or create a process or policy. We use it in everyday discussion about the content we put on our website, the ways we engage with our community, the products we build, and the internal actions we take.

TAGFEE started in 2007, when Moz's first investor, Michelle Goldberg, gave me a copy of *Good to Great*, Jim Collins's classic analysis of the elements that correlated with companies that achieved long-lasting greatness versus those that didn't. In his research, Collins identified seven characteristics highly correlated with companies that have remarkable financial and growth performance over long periods. One of these, which Collins called "First Who . . . Then What," posits that great organizations are made up of people who share fundamental core values and use these as their guiding light for decisions big and small.

Everything I've experienced as an entrepreneur, a CEO, an individual contributor, and a student of startups and business culture reinforces this concept: people who share core values and believe those values to be the most important part of their contribution to the world have the greatest potential to accomplish remarkable things together.

And conversely, when individuals in an organization don't align to the same values, every goal, project, and effort is undermined. Success, in any form, is made massively harder or easier depending on the degree to which your team, in the deepest part of their personal beliefs, shares an unwavering commitment to the same values and agrees with the implications of what values mean to an organization. This quote sums it up:

> The core values embodied in our credo might be a competitive advantage, but that is not why we have them. We have them because they define for us what we stand for, and we would hold them even if they became a competitive disadvantage in certain situations.
>
> **—Ralph Larsen, former CEO of Johnson & Johnson**

Moz, in its early days as a software company, latched onto this idea, and I believe we have always been at our best when we embraced and embodied our core values, and were often at our worst when we strayed from them.

Geraldine is TAGFEE's author. That might come as a surprise, but she was a copywriter by background (and is now the published author of *All Over the Place: Adventures in Travel, True Love, and Petty Theft*. You should probably get a couple of copies, in case you want to read it more than once). When we started down the path of identifying Moz's core values, she was an obvious choice to assist. She knew me incredibly well, she was familiar with the company and all eleven of our employees at the time, she was a talented writer, and she didn't charge very much (important to an early-stage startup).

Each of Moz's employees went through a written exercise, listing the traits and qualities we admired most in others and aspired to ourselves. We talked together about what we wanted to be as a company, what we wanted to stand for, and what we regretted from our pasts—personally and professionally. The notes from these exercises and discussions were passed on to Geraldine, who transformed them into a written document that we then shared, edited, commented on, and returned to her for a final version. It begins:

Moz's Guiding Principles

This document represents the rules we have created and ideals we strive towards for all the work we produce as a company. We embrace these as the embodiment of who we are, why we exist, and what we endeavor to achieve in every arena—from software to website content to actions in the workplace and on the road as representatives of Moz.

M 9780735213326_Lost_TX.indd 143 7/19/18 11:16 AM

And the six values themselves are:

- **Transparency**—We believe in sharing what we know, learn, and do with everyone who's interested. We reject secrecy, obscurity, and opacity in all its forms and strive instead to make the worlds of marketing, of SEO, of software startups, and of Moz itself open and accessible to all.
- **Authenticity**—We hate pretending to be people we're not or hiding our true identities, thoughts, or feelings in our work. We despise corporate, inauthentic behavior and the trappings of the business world that hold back our humanity or diversity. We will always work to make Moz a place where all of us can be our real selves.
- **Generosity**—We believe in giving back without asking for anything in return. Our goal, greater even than growth or financial success, is to help make our peers, coworkers, and the world of marketing a better, more nurturing, giving environment.
- **Fun**—The work we do can be challenging and stressful, but we believe work is only work if you make it so. We aim to make our jobs and the jobs of those around us enjoyable, rewarding, and humor filled.
- **Empathy**—Our most important value, empathy, demands that we put ourselves in the shoes of others and see things from their point of view. We endeavor to create products, content, interactions, and environments that are welcoming and respectful to all. We believe the best kind of empathy is that which aims for the most long-term good, not just a short-term veneer of niceness. Our goal is to apply this empathy with the highest priority to our community, audience, and customers, then to ourselves, and finally to our shareholders and investors.

- **The Exception**—If everyone else is doing something one way, we believe there's innate value in finding an alternative path. Moz strives to be unique, innovative, and weird. We cut against the grain and hope to stand out as an exception to the rule.

You can find the original version online in a blog post I published a little more than a year after its original creation.

It's very possible that TAGFEE's values resonate with you personally. But it's also completely okay if they don't. No two organizations should have precisely the same values, and values should not be arbitrarily created and then forced upon a team. Moz's values have worked for us because they come from a deeply personal place inside its founders' and early employees' beliefs, shaped by our experiences as people and professionals. They are not designed to work for everyone or appeal to everyone. Core values have proven their worth to me, as a founder, CEO, and employee, because they provide *three* powerful, unifying organizational forces:

The *first* is a shared commitment among the team. In a company with hundreds of people, there will naturally be tension, disagreement, and occasional discord. But core values help everyone know, from the day they're first interviewed, that unifying beliefs bind us together. Even when we disagree on how to accomplish our goals or on whether we have the right goals, we at least know that our deepest foundation is the same. My in-laws fight constantly about pretty much everything (usually in their native Italian, so I catch only every third or fourth word), but they all agree when it comes to politics, a rare thing in American families. That shared consensus of the most important parts of what's right and wrong for the world and its people can take a Thanksgiving dinner going off the rails and bring everyone back together. The same is true at a company with resonant, shared values.

The *second* is a set of blueprints for decision making. When we're faced with a challenging decision, lots of data, intuition, and analysis will naturally be a part of the process. But these can be powerfully bolstered with our values acting as guardrails we stay inside. As with the example from this chapter's opening of requiring a phone call to cancel, we are able to use our values to help determine how to invest in improving the company, its products, its people, and its growth. I'll share some more examples of this process later in this chapter and in the book.

The *third* is evaluation criteria for retrospection. Most companies do some kind of retrospecting (looking back at previous decisions, projects, and investments to determine whether they were worthwhile), and use inputs like return-on-investment, cost/benefit analysis, and other varieties of metrics depending on the circumstance. The addition of fit-with-core-values to these operations delivers a special kind of insight, and one that both helps reinforce those values and sets up future investments for better success. Consistency and commitment loom large in the human psyche, and when it's perceived by your internal team, external customers, and wider audience that your organization maintains them, your brand benefits.

We Hold These Core Values to Be Self-Evident

Values-driven organizations have proven that even in the early days, there is an unassailable magic that comes from having the right people on the team, and for those people to share a set of common values (alongside a shared mission and vision). The all-too-common problem comes when founders and leaders believe that recruiting can be separated from values alignment.

Collins explains this brilliantly in his essay on aligning action:

> In describing the alignment process, I have assumed that your organization's core values are already clearly defined—a big assumption. Let me make a few points about identifying core values, for without this stake firmly in the ground, there can be no effective alignment.
>
> First, you cannot "set" organizational values, you can only discover them. Nor can you "install" new core values into people. Core values are not something people "buy in" to. People must be predisposed to holding them. Executives often ask me, "How do we get people to share our core values?" You don't. Instead, the task is to *find* people who are already predisposed to sharing your core values. You must attract and then retain these people and let those who aren't predisposed to sharing your core values go elsewhere.
>
> **—Jim Collins, "Aligning Action and Values"**

I don't think there's an easier or more tempting mistake to make than to hire someone who has great ability or a great track record of performance with the recognition that while they're not a match with the current values and culture of the company, over time you believe you can bring them into alignment. I've made this same arrogant move multiple times over my career, and each time it ended in something between disappointment and disaster. The visual on page 148, originally from Moz's investor Brad Feld, sums it up.

In the professional world, we're accustomed to hiring for competence. It's been drilled into us by popular culture and long-held business practices that the goal of hiring is to recruit someone to the team

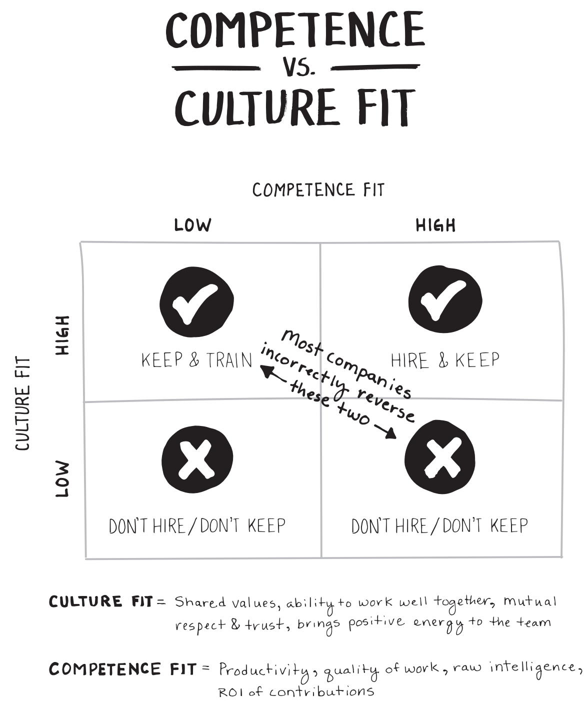

who has demonstrated skill and experience in a similar role. That's not a terrible thing to include in a hiring process, but if you want extraordinary results, it can't be the only thing you seek out.

Instead, you need a hiring process that considers core values and broad culture fit. Don't arrogantly presume you can transform a person who isn't predisposed to believe in or share your core values into someone who is. Build that screening into your interview process, your onboarding process, and the way in which contributions are judged. Use it to inform how raises, recognition, and promotions are given. That's the only way you'll prove to your team that values are on par with work output.

Values Demand Vigilance

It's oh so tempting to let a few values mismatches and clashes slide because an employee is doing high-quality work or has skills that seem hard to replace. But every time we've done this at Moz, it's backfired. Sometimes it's quick and the damage is minimal, but other times we've kept someone around far past the time they'd proven their values to conflict with our own; it wreaks havoc on morale.

Take Maya Angelou's advice: when someone shows you who they are, believe them the first time.

A few years ago, Moz hired a longtime veteran of a big software company. He came highly recommended with a number of leadership positions on successful products in his past. But early on, he clashed with a number of folks on the team. He played politics well, though, and neither his managers nor the broader Moz leadership felt his transgressions were untenable. Unfortunately, his peers and subordinates at the company, especially those not on his team, came to feel that Moz's commitment to TAGFEE and to upholding core values must be slipping. They assumed, given that we'd recently raised a large round of venture financing, that people like this guy were "the new normal," and it wasn't that leadership didn't know about his cultural mismatch, it was that they just didn't care as much anymore. So long as he performed, they figured, we'd keep him around. In their eyes, Moz lost some of its magic, and its leadership lost some of their credibility, especially on values issues.

I found out only after he left that in addition to these professional conflicts, he'd made sexist remarks and jokes to some of the young women on our team. He'd brought up inappropriate subjects at the office, sent around borderline offensive memes and links, and verbally bullied people who'd since left the company. I was heartbroken. I asked our head of HR why she'd never mentioned any of these things or taken action.

Her response shocked me: "You're the first person I've ever heard this from," she said. "A few people have complained about him, but never with these specifics."

His behavior had gone unreported. I asked a couple of the folks who told me about it why they'd never raised these events with their manager or with HR. Those replies were even more heartbreaking: "I didn't think it would do anything," and worse, "I figured everyone knew but didn't care so long as his projects did well."

Hearing stuff like that was like a knife in my gut. This company I'd built, that I'd worked so hard to make into a place that cared about people and values, had clearly changed. It wasn't because we'd hired dozens of vindictive or evil assholes. It was because letting even a few people (and this example I'm sharing sadly wasn't alone) break our core values repeatedly without visible action from leadership led to the normalization of discordant behavior.

There are *three* common ways values fail at organizations:

First is when they're viewed by the team as merely paper platitudes, hung on a plaque on the wall but not consistently enforced. If you overlook your stated values when they come into conflict with an employee who's performing well (or perceived by managers to be performing well), you're revealing reality—that performance (or perception and politicking) matters more than embraced values.

Second is when those values are created because the founders or executives thought they'd help build the cultlike environments that Silicon Valley culture broadcasts as a "must-have" element of startup recruiting and retention. You've almost certainly seen or heard about companies whose values have become the stuff of parody: values like "hustle," "work hard, play hard," "get shit done," and "always be shipping."

If any values you hold deeply are ones that, even if you found them to be a competitive disadvantage, you'd still uphold despite that conflict, then by all means, keep them. But if your company's values are merely marketing for potential hires, don't call them values. Be honest with yourself and your staff. Make them part of your recruiting materials and your internal lingo if you must, but don't attempt to pull the wool over your own eyes and everyone else's by calling them something they're not.

Real values have costs. They're difficult to embody. A lot of people (but hopefully not the ones you recruit) will disagree with them. People internally and externally should, at least some of the time, view them as a barrier to making a financially beneficial decision.

Real values are truths you hold to be more important than making money. They will come into conflict, and you'll have to make that hard decision and show everyone on your team why you're choosing that path, not just once, but over and over in order to instill the idea that values mean something at your organization. Because, usually, they don't. People who've been in the professional world for even a few years get pretty jaded about "company values." That means you have to go above and beyond to prove that you take them seriously.

If you're not willing to sacrifice and to make money-costing and painful decisions that bias to your values, don't bother having them. Say, instead, that your core values are financial growth and monetary success. You'll attract like-minded people who appreciate your honesty.

When you fail to live up to the expectations you've created with values statements, no amount of pretending can distract the incredibly talented, smart people that startups recruit from the uncomfortable truth that these so-called values can be willfully violated so long as the real goals (financial or otherwise) are met. Real values are proven through hard decisions and reinforced through recognition of times when they've caused real pain and still been "the right call." When your employees go out for drinks and bemoan the loss of revenue or opportunity and say, "Well, I guess they really believe in X more than just making money," that's when you know real core values are a part of your organization.

Third is when values aren't publicized at all and must be discovered over time through trial and error or by watching the reactions of scared employees in meetings with management. Working in environments like this is frustrating, tiring, and off-putting for those who can't quickly learn or keep up with the system. Every company has unspoken rules and idiosyncrasies that take time to learn. But values should be explicit, because they act like the operating system for a person's employment, affecting every action and decision they make. When you force people to figure out a secret, unwritten code for behavior, you are guaranteed to drive away a fair portion of otherwise talented contributors. Worse still, all the benefits you could achieve from recruiting and hiring with explicit values are forfeit. If you can't be bothered to identify, amplify, and reward your beliefs, there's no way for your recruiting process to seek out those shared beliefs in new employees.

A team with shared culture and shared values will, almost always, outperform a team without these elements. Why?

- **Retention**—It's far more difficult to retain team members who fundamentally disagree with how things should be done and why, even if they agree on what work to do. And constantly rebuilding a team through hiring is both exhausting and inefficient—people do their best work a year or two into working together with the same group. They learn one another's intricacies and idiosyncrasies, anticipate needs, establish efficient communication and process patterns, and know what works and doesn't for their teammates. According to Namely's analysis of more than twenty thousand startups in the United States, average employee tenure is only 10.8 months. 10.8 months! If you can improve yours, you'll have a powerful competitive advantage.

- **Motivation**—The difference between working with people you like, trust, and agree with versus working with those whose tactics, style, and ethics you question is immense. The former promotes the assumption of good intent (among the most critical elements for team bonding, productivity, and quality of output). The latter fosters political environments where work quality suffers.
- **Cohesion**—It is massively easier to ask a team to commit to a road map, a project, or a process (even when they disagree) if core values and culture already bind you together. When those elements are lacking, so, too, is the basic structure for what makes a compelling argument or what elements should be part of decision making or which paths are available and worthy of consideration.

Homogeneity Hobbles Innovation

There is one big, fatal flaw that often accompanies the pursuit of shared culture and values: uniformity.

Startups and early-stage ventures need diversity. I mean that both in the sociological sense (i.e., not just young, white men from the same country and background) and in the broader, thought-pattern and experience sense (i.e., everyone in the company thinking the same way and bringing only a single set of professional/personal experiences). These two might seem at odds: diversity versus shared culture. They are not.

In fact, shared culture combined with diversity should be exactly

M 9780735213326_Lost_TX.indd 153 7/19/18 11:16 AM

the combination you're seeking in a team. Diversity, culture, and values might sound like a complex, tough combination. I agree it's challenging to build a team with these attributes, especially if your own background lacks exposure to diverse people and/or those who share your values. But real magic happens when these come together.

A major roadblock is the easy-to-make assumption that diversity of employees, especially early on, isn't desirable, because it conflicts with or contradicts the goals of shared culture. When I hear this pushback from founders, it gets my blood boiling. This is a fundamental misinterpretation of what culture, values, and diversity mean. When we talk about diversity in a founding team or an early-stage organization or a larger company, we're talking about recruiting and hiring a wide range of people from different backgrounds, ethnicities, ages, genders, and identities (in all their forms). People with these varied attributes are not fundamentally at odds with your culture or your values. Women of color and disabled veterans and Asian men in their sixties and young people with nonbinary gender identification can all share the same core values and believe the same things about how a company should be structured and how people can work together.

I'm not saying that every randomly assembled group of diverse people (or every randomly assembled group of straight, white, cisgendered men in their thirties from New York) *will* share these beliefs. I'm saying that recruiting diverse people who do share these beliefs is a massive advantage—a cheat code—for your organization.

Diversity is immensely desirable because it improves perspective, empathy, and creativity. A diverse group brings unique life experiences and, as a result, a unique ability to contribute that non-diverse groups can't match. It's often hard to know when you're benefiting from diversity (or when you're being hurt by lack of it) because the input given by people is seldom directly attributable to their background, but I'll give six short examples.

1. When we started recruiting to grow the company, gender diversity made a big difference. My mom was my cofounder. Sarah was my COO. Our board comprised more women than men until 2012. That meant Moz was seen in Seattle as a place where women in the startup world were welcomed and encouraged. I hired a woman (Kate Matsudaira) as our vice president of engineering in 2009. She was a critical hire who seriously upgraded our engineering practices. I suspect I never could have swayed her to join an all-dude company. The same held true for many other folks we brought aboard—junior and senior, technical and not. Imagine excluding 50 percent (or more) of your potential workforce because you recruited only your same-sex drinking buddies. Not a smart move given how hard it is to hire in our world.
2. In developing personas for our products, we started with what can only be called the whitest, most generic naming conventions ever. Thankfully, some more thoughtful and diverse Mozzers noticed this convention, and noted the subtle effects that subtle biases like names and genders, even in fictional characters (like personas) can have on how we perceive the world of our customers and design for them. If a persona is called "College-grad Chad," there's a certain associated gender and background identity for most folks. Our designers assumed Chad could read small text. Our engineers assumed Chad was familiar with advanced query modifiers. Our marketers assumed Chad was active on Twitter and Instagram. By moving to more inclusive persona-naming conventions and descriptions, we brought to the process a more accurate view of the diversity of the customers we were actually designing for,

building more accessible products and marketing them more thoughtfully in the right places and ways.

3. At one of Moz's executive team lunches some years back, someone brought up the City of Seattle's decision to stop using the term "brown bag" to describe lunchtime presentations. I thought it was very odd until our CTO at the time, Anthony, who's black, noted that the phrase was used to segment and classify people based on the color of their skin, and he'd personally been subjected to it as a kid in eastern Washington. We immediately got rid of the term at Moz and switched to "lunch and learn" (which has the added benefit of being much more understandable to our non-American-born employees and guests).
4. At another executive meeting, we were all talking about the various "tribes" at Moz (a term we'd been using to denote groups across teams that work together on projects). Annette, our CMO, whose background is American Indian, wondered if we could choose another word to describe these groups. . . . Cue another head-smacking minute, and a shift in terminology.
5. When Sarah (now our CEO, at the time COO) was pregnant with her son, she noticed the lack of private rooms at Moz where moms and moms-to-be could comfortably take care of themselves or their kids (and that restroom stalls were not an option). That pain and awareness caused us to correct the issue with rooms specifically for that purpose, but if she hadn't been part of the office design team, it's very possible we wouldn't have been sensitive to the issue.
6. Finally, and perhaps most broadly, in a design review of the first version of Moz Analytics, one of our software tools, Sarah, and a number of other women on our teams who looked at the product, took exception to several elements of the design (color, layout, font

> usage, and word choices). When we modified these in response to their input, metrics improved. David Mihm (whose company we acquired in 2012, and who is partially color-blind) helped identify which contrast variations he literally couldn't see. Martin York, one of our senior engineers, who has dyslexia, commented on form inputs missing auto-corrections for common mistypings. One of the older members of our team noted that she had difficulty in parsing the text because the lines were too close together (a.k.a. overly tight leading). I couldn't have caught these issues myself and neither could a team of people with my background, gender, age, or abilities. Diversity made our product more accessible, improved engagement, and reduced customer frustration.

This is why diversity is so well correlated with success everywhere from early-stage startups to Fortune 500 boards of directors. Research from McKinsey showed that more gender-diverse companies outperform their less gender diverse peers by 15 percent while more racially diverse teams outperform their peers by 35 percent. PE Hub, *Venture Capital Journal*, and Women VC published a joint report analyzing the returns of investment funds with gender-diverse versus mostly male and all-male investor teams and found the more balanced teams had returns 3.78 times higher than their less-balanced peers. When First Round Capital analyzed the performance of the hundreds of investments in its own portfolio, it found that teams with at least one woman founder performed 63 percent better than all-male founding teams.

If I were founding a new startup today, I'd do almost anything to have a feature that's 63 percent higher correlated with better performance.

Where founders need to seek out shared attributes isn't in what we look like or where we're from; instead, focus on ethical beliefs and the right ways to operate a company. You want a team that has significant overlap in its answers to questions like:

- What traits and behaviors should be rewarded and recognized in our employees?
- Which ones should be discouraged?
- What criteria should be applied to people we hire, those we promote, and those we let go?
- What makes someone a good person versus a bad person?
- How should hard-to-resolve conflicts be handled at our organization?
- What's your preferred form of communication and why?
- What enables you to deliver your best work? What stops you from it?

You've probably seen stories in the tech press about how a startup is asking potential hires about *Star Wars* versus *Star Trek* or which craft beers they prefer or what sports they watch in order to help determine "culture fit." These types of questions are awful not only because they promote homogeneity of thinking ("we're all from the same town and love the same soccer team and the same video games" is *not* the kind of shared culture that will help your organization succeed), but because they, intentionally or not, bias to people with homogeneity of experience, too.

Twentysomething white men from upper-middle-class and wealthy upbringings are the most plentiful founders in the technology world. I fit into this group myself. My dad was an engineer at Boeing. My mom was a designer, a marketer, and an owner of her own small business. They're both ethnically Jewish, but not practicing.* My parents made decent money and were extremely thrifty, enabling them to pay for college for all three of their kids (though only one of us, my sister, actually graduated). We grew up in a rural area outside Seattle. And I get how tempting it is to find someone like me as a cofounder for a new venture. I already know lots of other white and Asian men in the Seattle area who grew up in middle-class families, are in their midthirties, don't have kids, like computer games, and cheer for the Seahawks.

What's wrong with starting a company with friends I already have, whom I know I get along with, and whose interests and passions match my own?

It's not additive.

Together, my peers and I of similar background and identity combine for only a very slight bit of extra perspective. Try as we might, our outlooks and our framing of events will be colored by who we are, where we've lived, what we've experienced—the benefits of diversity will be generally lost to us. And, as a result, our ability to empathize with, design for, market to, and serve broader groups with whatever we make will be limited. Additionally, we'll almost certainly find it harder to recruit diverse early employees, which will have a domino effect on our hiring and team composition for the long term. If you wonder why so many startups are so homogenous, or why so many serve only to address the problems of a very thin slice of the world's populace, look no further. It's usually not intentional, evil,

*Interestingly enough, in the 1970s, Jews successfully lobbied the US government to be classified as "white," and have mostly fit under this umbrella since (well . . . at least until the 2016 election).

biased hiring practices, but rather deep-rooted, systemic drivers like the design of interview questions and where job postings are placed and the practice of recruiting through friends and family networks.

Moz itself struggled massively with diversity, especially on the technical side. I think this is because we unintentionally fell prey to all of these inherent biases. We looked for familiar experiences. We sought out many of our early employees through the people we already knew. We made no intentional efforts to create a diverse pool of candidates from which we could hire. Like many companies, we didn't think much about diversity for years until, one day, we did, and it looked awful. In 2012, more than 90 percent of all of Moz's engineering hires were white or Asian men in their twenties and thirties. We could have been the poster child for stereotypical tech monocultures. And this was despite having a woman (Kate) and then a black man (Anthony) serve as CTO.

We didn't just struggle on the diversity front, though; we struggled on shared culture and values, too. We'd hire people for their experience, skills, and ability without a consistent, intentional process to determine whether they supported TAGFEE or agreed with how we did things at Moz.

We addressed both issues in the space of a year. First, we built a new process called the "TAGFEE screen," wherein members of a team not making the hire would spend time with candidates during an interview day and have discussions designed to elicit alignment with our values and our culture. For example, if our Big Data engineering team was interviewing a candidate, that person might go out to lunch with two folks from customer support to talk about values and culture issues. If the customer support interviewers had substantial concerns about the engineer's proclivity toward (or away from) TAGFEE and Moz's people-centric culture, they could veto that hire . . . even if the engineering team thought the candidate was a perfect skills match.

Next, we looked into our practices around diversity. We found, not surprisingly, that a high proportion of our candidates were sourced internally through referrals. And of course, few of those candidates had any diversity of background or experience. They went to the same schools, lived in the same places, had the same sorts of upbringings, etc., as our existing team. We made a number of intentional efforts to change by investing in and supporting local programs aimed at promoting diversity. These included Returnship, a program for parents (usually moms) returning to work after a few years off for childcare; Ada Developers Academy, a program for women learning to code that we hosted at our offices; and TAF Academy, a school program to help get kids from underprivileged backgrounds access to science, technology, engineering, and mathematics education. Our goal was to gain exposure to potential candidates whom we'd otherwise miss.

The results were remarkable. A couple of years after implementing, our engineering organization has improved gender-balance and background diversity, more than tripling the number of people of color and women engineers on the teams (though we still have a long way to go). Other parts of the company—customer success, finance, marketing, product, and facilities—have become more diverse, too. We've discovered places to post jobs we never knew about, language that was holding back diverse candidates from applying, and new networks of people to help us attract broader groups of applicants.

Pro Tip: We use and love Textio (https://textio.com/) to analyze our job postings, about page, and other recruiting-focused content to make sure we're using inclusive, unbiased language. It's worth checking out and has a number of free and low-cost options.

On the shared values and culture front, we've shown similar improvement. From 2011 to 2013, we hired very fast. This brought a cohort of Moz employees who reported less job satisfaction and had higher voluntary turnover than prior cohorts (i.e., those folks were more likely to leave the company in a shorter period of time). After the implementation of our TAGFEE screens and our more considered focus on culture and values in the interview and hiring process, we improved voluntary retention. We have, in my opinion, also improved the caliber of new employees joining the team. Teams work better together, get more done, and produce higher-quality results than ever before.

If you're building a team from scratch or are in the early stages of recruiting and hiring, I hope you'll learn from our mistakes and from the data. Hire people who share a belief of what deserves a promotion and a raise versus a reprimand or coaching. Hire people who are naturally inclined toward your values and who want a place that's willing to sacrifice short-term growth or financial success in exchange for adherence to them. But don't hire only the people who look like you and see the world through the same set of experiences. Bolster your potential for high performance and for broader customer empathy by intentionally seeking out diversity. These two elements, when combined, forge the underpinnings of a remarkable team.

CHAPTER 11

LIVING THE LIVES OF YOUR CUSTOMERS AND THEIR INFLUENCERS IS A STARTUP CHEAT CODE

A great way to build software is to start out by solving your own problems. You'll be the target audience and you'll know what's important and what's not. That gives you a great head start on delivering a breakout product.

—Jason Fried, March 2006

I thought I had an amazing idea. I thought it was going to change the world of marketing. I thought we were going to build a set of software that every business needed. I thought it would catapult Moz's growth and revenues. I thought wrong.

It was 2011 when Adam Feldstein, then Moz's chief product officer, and I sat down to plan out a project we called "Moz Analytics." The new product stemmed from a theory I had that in the near future, the siloed practices of social media marketing, search engine optimization, content marketing, public relations, and online brand marketing would all merge into a single set of tactics undertaken by the same person or group at an organization. I saw how social media and content marketing worked together to bolster each other. I wrote and talked at conferences about how many PR and brand-building efforts were merging with SEO. I saw how a few organizations had already combined these practices into remarkable flywheels that generated returns greater than the sum of their parts. I knew these practitioners would need tools that worked together to optimize their efforts, track their progress, and compare themselves against the competition.

Notice how I never thought to validate my idea externally? How nearly every sentence in the passage begins with "I"? You can almost picture the train wreck on the horizon.

The Case of the Disappearing Conversions

In November 2013, after more than two years of planning and development, we finally released Moz Analytics. More than ninety thousand people had seen a preview of the product and signed up to get notified of the launch. It was by far our most successful product marketing campaign, and the buzz from our community felt extraordinary. Every day, we found new, speculative discussions about what might be in the product or how it might help with their work. It felt like we were releasing a blockbuster movie rather than a business tool for marketing professionals.

But internally, we were struggling. The delivery date slipped five different times. We'd had to replace our engineering lead on the product. Features were cut, then whole sections of the product were cut. The version we released was buggy and incomplete. But worse than that, it wasn't what our customers wanted. That month, I delivered an all-hands presentation (my last one as CEO) to the Moz team, detailing what had gone wrong and where we stood. This slide was the setup:

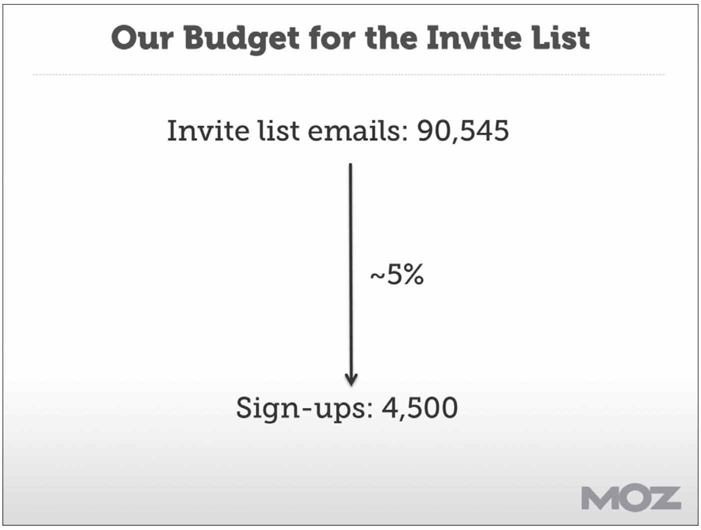

And the next one was the kicker:

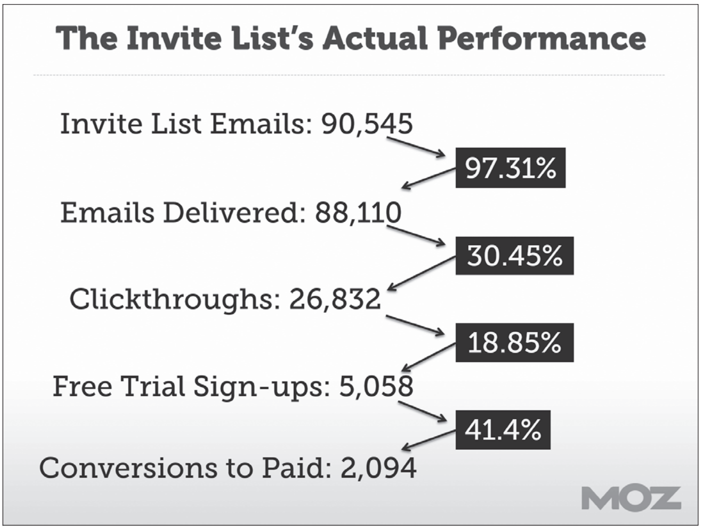

Of the 90,545 folks who'd said they were interested in our new product, only 2.3 percent actually paid for at least a month of the service. Worse still, those customers tended to quit the subscription far earlier than customers of our prior product. It was almost a year of additional work before the new product was performing as well for our customers or our bottom line as our previous one.

What did we get wrong?

First, we built software in the worst possible way. Rather than creating a small kernel of the eventual finished product, then iterating and adding until we had a satisfyingly useful and high-quality deliverable, we designed a massive scope of work and asked large groups of engineers (both internal and contracted) to work on separate pieces that would then (supposedly) fit together. Multiple contractor groups failed to deliver what we needed. So, too, did multiple internal engineering groups. Everyone was demoralized. Dates were moved back by months at a time. We had no ability to show off any version of the tool set until just a few weeks before the final launch, costing us invaluable customer feedback and time to make things better or rethink the product.

Second, because of our delays, we felt insurmountable pressure to release as soon as possible. Even though customer testing a few weeks before launch revealed loads of bugs and dissatisfaction, our morale as a company and leadership team was so low that we were desperate for a release. We figured we could launch with a "good-enough" product, then iterate until it was great.

But guess what? People judge by first impressions. When those 26,832 people visited the page announcing Moz Analytics and showing off what it could do, most of them disappeared, never to return. Many who tried the product came away unimpressed. The "word on the street" (or in our case, the web forums, conference halls, and social media discussions) said Moz had a crappy new product that wasn't worth the money. That reputation dogged us for three, long, growth-stunted years.

Third, we spent years building a product based entirely around a theory—one that proved to ultimately be false. Five years after my prediction that social media marketing, SEO, content strategy, PR, and all those other "earned media" practices would combine and be the responsibility of a single person/team, online marketing remains as specialization-heavy as ever. I was so confident in my knowledge of the field and my ability to predict what people needed that I failed to do the real research required to validate those assumptions. Instead of spending time with my customers and potential customers, I spent it with my product designers and engineers, dreaming up wild new things we could build.

It wasn't until October 2013, just a few weeks before we released that doomed product, that I did what I should have done years before: put myself in our customers' shoes.

If you're ever tasked with a large software project, learn from our mistake. Pare back your design until it's the smallest possible element of what you eventually hope to have. Show that to people you trust and get their feedback. Iterate on the fundamentals. Then build it one element at a time. Add functionality, data, features, visual elements, etc., until you've got something new to show your trusted advisers and beta customers. But don't release it broadly until the buzz you're getting from these groups is firmly in the "we love this and can't live without it" camp.

That Time We Reenacted the '80s Classic *Trading Places*

Wil Reynolds is the founder of SEER Interactive, a 150-person web-marketing agency in Philadelphia with a stellar reputation, ten years of steady growth, and a vast array of impressive client work. Wil and

M 9780735213326_Lost_TX.indd 167 7/19/18 11:16 AM

I have been friends for many years, having spent time together at numerous events and finding our way to more than a few dinners and after-hours gatherings. We both have strong opinions, a desire to make the world of marketing a better place, and experience building companies from the ground up.

Somehow, one night in 2012, at a drinking hole in Philadelphia, we'd shared a few whiskies and decided that the following October, we'd each take a week off from our own lives to inhabit the other's. It's the kind of conversation that usually never makes it out of the pub. We'll answer each other's emails! We'll live at each other's houses! We won't just be CEO in name only—we'll have real, decision-making authority! When we woke up and sobered up, it somehow still seemed like a good idea, so we made it a reality.

On Friday, October 4, I flew to Philadelphia with Geraldine. We took a cab to Wil's house in Northern Liberties. I learned how to care for his dog, Coltrane. Wil learned that, because we traveled so much, we'd already killed all our houseplants. We swapped email logins and critical passwords. We walked each other through the major projects and meetings on our schedules. We traded house keys. And on Saturday, Wil flew to Seattle and moved into our Capitol Hill apartment.

That week was indescribably challenging, intense, and rewarding.

Managing another human being's email, by itself, was a massive undertaking. I had to research new people, learn about projects, reach out to Wil's coworkers to ask for help with context, and constantly flex my best judgment muscles. I replied to emails from Wil's mom (with whom, delightfully, I'm still in occasional email contact to this day). I scheduled calls with new potential clients. I responded to team members and existing SEER clients who'd been instructed by Wil to treat his email as though it were any other week.

Wil starts his days early. I'm a night owl. And because of the time zones, Wil's schedule was already three hours ahead of mine. Getting to the office at the equivalent of five a.m. Seattle time took its toll, but I muddled through. Wil's assistant, Stephanie, was both immensely kind and unflinchingly strict. She had me scheduled most of the days, back-to-back, with meetings, events, or one-on-ones of one kind or another. My amazing assistant, Nicci, did the same for Wil. We both still rave about each other's EAs. (*Pro Tip*: If you can afford an exec admin, get to it! You'll boost your productivity threefold, I promise.)

I got mentored by Wil's team members in how they did many of their client projects and how they used software, including Moz's tools and several of our competitors', in their processes. With help from SEER's CFO, Larry Waddell, I dug into how they attracted customers and how they managed the complexities of waxing and waning client demand. I learned about SEER's finances and their management structure, about their promotion criteria and their team's strengths and weaknesses.

I even had an unexpected meeting crop up on Thursday, during which one of Wil's employees turned in her two weeks' notice. Remarkably enough, the same thing happened to Wil in Seattle as one of Moz's senior engineers announced his intention to leave.

We didn't just do each other's jobs; we lived the other's life. Geraldine and I volunteered at the Ronald McDonald House in Philly (a commitment Wil had made months earlier). I attended his monthly meeting with Covenant House, a charity that helps homeless youth get housing and support services. We ate dinners with Wil's wife, Nora (who couldn't participate in the swap due to work). I fed Coltrane and took her to the office. On Monday, everything felt awkward, but by Friday, both of us and our respective teams had reached a surprising level of comfort and normalcy with the swap.

I sat in on meetings with Google's advertising team and pitched

SEER's services to a pair of potential clients (they closed one deal but lost the other). Several SEER employees took me under their wings, showing me the processes they'd built and explaining the ins and outs of their site audits, their keyword research processes, how they managed client ad spend accounts, and more.

I watched SEER's consultants use tools. Sometimes it was Moz's, sometimes our competitors'. Two things stood out.

First, it was clear that every data point needed validation. If Moz's tool said there was a link from site A to site B, or that page X had a particular problem, the consultant would go check, manually, that there was indeed a link from A to B, or that the problem on X was the one we reported. Only after they confirmed the tool reports with spot-checks would they start to trust the results and assume accuracy. And if they found discrepancies, they'd switch to another tool, or worse, to an entirely manual process.

Second, in every case where SEER consultants were using tools, they'd happily switch from one hyper-specialized solution to another. If Moz's keyword data wasn't ideal, they'd use SEMRush or Übersuggest or Google AdWords. If our link data wasn't comprehensive, they'd move to Ahrefs or Majestic. There seemed to be no loyalty between tool providers and no barrier or switching costs from having to change up the UI, log in to a different tool, pay another monthly fee, or get data in another format.

The benefits I thought an all-in-one tool set created were being dismantled before my eyes.

That exposure to Wil's professional life and to SEER's inner workings was transformative. What started as a bit of a gimmick emerged as possibly the most intense, high-speed course in customer empathy imaginable. Instead of merely watching customers use our products or interviewing them about their work, I was living their lives, eating, sleeping, and breathing the challenges of an SEO agency. That week impressed upon me how different reality was from my expectations, and, along with the poor reception of Moz Analytics, made me question everything I'd assumed I knew about the web-marketing industry.

I desperately needed that humbling. I only wish it could have come sooner. It was too late to change the scope of launch and impossible to go back in time to warn my colleagues and myself that we were embarking on a voyage without the necessary compass. But that week in Wil's shoes made me more thoughtful about every plan and product I've committed to since. No longer do I trust solely my own judgment. I validate not just with interviews and testing, but by forcing myself to do the work, by using our competition's products to accomplish those tasks, and by spending time with my industry compatriots, often while on the road for conferences and events, to share side-by-side, working time with them on real projects.

Another big change I made was a return to the consulting world. Though I didn't charge for my assistance, exposure to SEER made me realize how crucial it was that I didn't just theorize and prognosticate, but actually got my hands dirty. Most of the projects I take on now are either nonprofits or via personal and professional friends I'm seeking to help. Often a half-day session with a team of marketers at a startup will uncover a huge blind spot in our tools or in the universe of available tools, and I'll return to Moz with a half-baked idea I can validate over and over before deciding if it was something real, and worthy of pursuit, or just a one-time problem.

If Life Swapping Isn't an Option . . .

Many paths exist to live the lives of your customers, but before you can do any of them, you need to know these people as people, not as "personas" or "sales targets." I've found repeatedly that when our product and engineering teams comb through user interview or review data, we tend to create features and products that are barely better than the already-established processes our customers are using.

I think this happens because of how we've been socialized and trained to use data in professional environments. We see numbers, we analyze them, and we make decisions based on how those numbers lean.

Say you're tasked with making software to help people manage their finances. You survey a large group of people about their spending habits, what they want to track, what information they need to be informed about, and where they have financial analysis pain today. The surveys and interviews reveal the ten most important categories of spending and that the proportion of spending over time and the absolute amount spent are both important. Thus, you create the same app that nearly every major bank and credit card has today to serve this purpose.

But if you knew the people personally, and spent time with them while they did their banking, financial management, and planning tasks, you might find that post-spending analysis isn't nearly as important or helpful as alerts before they spend, tracking progress against goals, or incentivizing healthy spending and saving behavior.*

*Yes, I realize that these types of features might dissuade a person from spending as much, running counter to the financial institution's goals of getting them to spend more (and hence won't be popularized in their feature sets).

Facebook had a great example of this several years back. Internal data showed Facebook that many users loved to reminisce by viewing photos from years past. The high engagement rate these photos received prompted the social network to launch a new feature called "your year in review." The feature was beloved by many, but also made headlines in the press because for some users who'd had tragic events, it was an unexpected, heartbreaking, and painful experience. Eric Meyer famously wrote a blog post called "Inadvertent Algorithmic Cruelty" about the experience of losing his daughter, then seeing Facebook push photos and stories of her to him at unexpected times.

The data was clear, but empathy for the real-life experiences of a crucial subset of Facebook's customers was missing. I want to believe that if an engineer on the Facebook year-in-review product had shared experiences like Eric's, or if they'd known people who did, that feature would have launched with some preventative logic built in for those who'd experienced tragedy like his.

How do you get to an empathetic place for product design and development? Create regular customer exposure for your team and yourself.

That exposure can come through extreme efforts like my CEO swap with Wil, or it can come from more subtle actions. There's a few that have worked particularly well for Moz's product folks and for other startups I've worked with, including:

- **Conferences and Events:** I get a lot of marketing value from speaking at events, but equally valuable is the exposure to professionals in our field who need and use the tools we offer. Hallway conversations, session Q&As, coffee meetings, and after-hours hangouts offer a wide range of experiential cases from which you can gain perspective and insight. Just be sure to have a few consistent, open-ended questions that get to the core of your customer-empathy issues.

- **Volunteering/Apprenticeships/Internships:** A handful of startup founders and startup product owners have taken the innovative step of volunteering a day, a week, or even a few months through an apprentice/internship (formally or informally) with their customers in order to learn what their day-to-day work, challenges, and current solutions look like. If you're in the early stages and have the ability to *be* the customer you're going to serve, even for a very limited period of time, I highly recommend it.
- **Paid or Pro Bono Consulting:** This is how Moz got its start. We were consultants first, built software that we ourselves needed, and then opened it up to a broader audience we'd built via our blog. Nowadays, I do this through pro bono consulting, and several of Moz's other product contributors still do some independent paid consulting. It's true that consulting doesn't always provide perfect insight into the issues or work faced by your target customers (unless they, too, are consultants), but it can contribute experiences and build relationships you otherwise couldn't get.
- **Teaching:** Not only does teaching require you to understand a subject or process deeply, it also gives you exposure to a wide range of practitioners or future practitioners of your subject matter. Those relationships bolster empathy as you show folks the what and how behind a process. It's no surprise that so many professors and educators are recruited as startup advisers.
- **Hiring or Contracting:** If you or your current team have no bandwidth or no passion for embedding yourselves into your customer processes, there's no shame in recruiting to help fill this gap. We've done this many times at Moz, hiring SEO professionals who know the field well and have used dozens of tools for years to assist us in

> building better software for customers like them. The key is identifying those individuals who can translate their own problems into more global solutions and have a product-driven mind-set, rather than focusing exclusively on their own processes. We've done best with this through the worlds of social media and blogging, ID'ing folks whose public contributions to the field clearly show an affinity for holistic, empathetic pattern matching and helping the industry as a whole.

Thankfully, after learning from my mistakes with Moz Analytics, a couple of years later I was given a second chance to build a new product. In this next chapter, I'll walk you through how we did it and try to reverse engineer why it worked so well.

CHAPTER 12

GREAT PRODUCTS ARE RARELY "MINIMALLY VIABLE"

If you are not embarrassed by the first version of your product, you've launched too late.

—Reid Hoffman, March 2011

In the past decade, the "lean startup" movement has had, arguably, a greater impact on the approach that product designers, engineers, and entrepreneurs take than any other. The fundamental concept is undeniably compelling: craft a version of your product (and company) that requires the least amount of time and effort to validate whether the problem you're solving is an important one that real customers will pay for or use. This "minimum viable product" will help you learn faster, iterate faster, and survive longer on less money. It's a powerful way to overcome many of the problems that plague (and often prematurely kill) young companies and new product efforts.

But it also leads people to create a lot of crappy, barely useful products.

M 9780735213326_Lost_TX.indd 177 7/19/18 11:16 AM

We'll Learn a Lot Once the MVP's Out

In late 2014 and early 2015, I worked with Moz's big data and data science teams designing a minimum viable product (MVP) to help people identify websites that Google might consider to be spam.

We started with a few assumptions about the field of web spam and SEO that we then validated through research and customer interviews:

- Getting links from spam sites can potentially harm rankings and visibility in Google.
- Knowing which sites are spam is hard because Google won't label them (if they did, spammers could easily see what passes Google's filters and what doesn't).
- If you take the time to perform searches in Google related to a website, and see that it doesn't rank for any terms and phrases that it obviously should (e.g., if Moz.com didn't rank on the first page in Google for "Moz" or "Moz com," we'd know something was wrong), there's a good chance Google's penalized or banned that site for spam.
- If spam does link to you, Google recommends using the "disavow tool" system in Google Search Console, but you must be incredibly cautious, because disavowing non-spam links to your site can result in massive traffic losses. (Cyrus Shepard, Moz's head of SEO at the time, tested this by disavowing all the links to a site he owned, and it plummeted in the rankings . . . lesson learned).*

---

* You can see more about this humorously fatal experiment here: https://moz.com/blog/google-disavow-tool.

- Many of our interview subjects said that fear of Google penalties and the constant need to identify and validate spam versus non-spam links was driving them up the wall and consuming a lot of their SEO work time.

Based on these learnings (and others I've excluded for brevity), we decided to build a "Spam Score" into our web index. This score would help indicate the degree to which a site might be perceived as spam by Google, and thus a potentially risky place from which to get or have links.

The MVP process used a clever bit of research from our head of data science Matt Peters. Long story short, Matt and I dreamed up nearly a hundred potential factors that might be correlated with sites Google had penalized or banned. We then generated a large list of websites that didn't rank for their own brand or domain name (indicating they had been flagged by Google) and looked at the relative connections between all those hundred factors and the penalized or banned websites. In the end, we found seventeen factors that were relatively good predictors of whether a site had drawn Google's ire.

We called them "spam flags" and saw in our research that the more flags a website had, the more likely it was to be penalized in Google's rankings. The flags included things like the length of the domain name (turns out spammers often have very long domain names) or the presence of many external links with very little content. Having a few flags wasn't a particularly bad thing—most websites triggered at least two or three. But if a website triggered eight or more of the seventeen flags, it was more likely than not to be penalized.

The great part about Spam Score, at least for us at Moz, was that it required a relatively small amount of additional work (all in, about three months of effort from five people, though it spread out across almost a year due to overlapping priorities) to include in our data sets and to publish in our tools. We knew that initially it would receive some fair and justified criticism, and we expected folks to have concerns like:

- Moz's web index wasn't large enough (at the time) to cover all of the domains that may be spam, and thus couldn't provide a comprehensive list of sites to disavow with Google.
- The percent-risk model can be confusing. Many people would prefer a model that simply showed whether or not a domain was penalized by Google (rather than a percent-chance tied to a count of features), but we didn't have the bandwidth to make that happen.
- Spam flags could be misconstrued as a potential problem for one's own website rather than a filter system for reviewing links from other spammy websites.
- The scores of five through eleven (out of seventeen) could be particularly vexing because they indicate a higher risk of penalization, but could also be totally innocuous.
- The flags weren't actually the spam signals Google uses (we don't know what those are because Google doesn't disclose them). They're simply well correlated with sites that have drawn penalties in the search engine.

At launch, we figured, despite these issues, our MVP would still help a lot of people and, like all good MVPs, it would help us learn more about our customers and what they wanted from a spam-identifying product long term.

But here's the kicker: Our research had already revealed what customers wanted. They wanted a web index that included all the sites Google crawled and indexed, so it would be comprehensive enough to spot all the potential risky links. They wanted a score that would definitively say whether a site had been penalized by Google. And they wanted an easy way of knowing which of those spammy sites linked to them (or any other site on the web) so they could easily take that list and either avoid links from it or export and upload it to Google Search Console through a disavow file to prevent Google from penalizing them.

That would be an *exceptional* product.

But we didn't have the focus or the bandwidth to build the exceptional product, so we launched an MVP, hoping to learn and iterate. We figured that something to help our customers and community was better than nothing.

I think that's my biggest lesson from the many times I've launched MVPs over my career. Sometimes, something is better than nothing. Surprisingly often, it's not.

Spam Score launched on March 30, 2015, and while we did receive a good bit of positive feedback, we also got a lot of criticism, confusion, and questions. The score's design was suboptimal. The way the flags aligned to a percentage risk model wasn't intuitive. Many users focused on the flag count for their own website rather than the flags of the incoming links to their sites. These were things we knew would happen in the design and construction phase but pushed to the back burner in favor of a faster release.

Marie Haynes, one of the world's foremost experts in the field of web spam and Google penalty issues, left a comment in the launch blog post that summed up a lot of the sentiment around the release:

> *I wanted to like this tool, but I am really concerned that it could do more harm than good. Perhaps I have misunderstood*

its purpose. If used as an adjunct to a manual link audit, it could be helpful. But to me, it came across as an all in one solution to link problems. I think that other people are going to assume this as well.

We'd talked to Marie during the development of the metric. We knew her concerns. We knew she was massively influential in the space and that her approval and support (and others like her) were a great barometer for our success at solving the problem, but we chose to launch while we were still "embarrassed" by our first version of the product, rather than waiting until we could develop something better. Perfect is the enemy of done, right?

Six months after launch, while looking at our product performance metrics, we noted that spam score had become mildly popular with a small group of our customers (about 5 percent of the folks who regularly used Open Site Explorer visited the spam score section), but it had no observable impact on free trials, on vesting rate, on retention, or on growth of the Moz Pro subscription overall. In other words, we'd probably have seen exactly the same performance in our customer base and growth rate if we'd never launched Spam Score.

Great use of (at least) $500,000 in data collection, research, and engineering time, eh? Thank god I'm the founder . . . otherwise I might have been shown the door.

Do MVPs Have to Be So Minimally Viable?

The problem with MVPs, and with the "something > nothing" model, is that if you launch to a large customer base or a broad community, you build brand association with that first version. To expect your initial users (who are often the most influential, early-adopter types you'll attract—the same ones who'll amplify the message about what you've put out to everyone else in your field) to perceive an MVP as an MVP is unrealistic.

In my experience, our customers (and potential customers) don't see new things and think: "Oh, this must be their initial stab, and while it's not exactly what I want or need, I can see that it's a product I should pay attention to and help support, because eventually I can imagine it getting to the place where it really is useful and helpful to me."

Instead, they (usually) see new things and think: "Is this interesting? Does it do what I need? Is it way better than what I already use? Is it worth the hassle of learning something new and switching away from what I've always done?" and if the answer to those questions is a "no," or even "Well, maybe, but I'm not quite sure," your product is unlikely to have substantive impact.

Worse, I've found that when we launch MVPs, the broad community of marketers and SEOs who follow Moz perceive our quality to be shoddy and our products to be inferior. I've termed this brand reputation that follows an initially incomplete, minimally viable product's launch the "MVP hangover." It seems to follow the product and even the broader brand around for years, long after we've iterated and improved to make the product truly exceptional and best-in-class.

My theory about MVPs applies differently to different stages of your organization, based mostly on reach:

For an early-stage company with little risk of brand damage and a relatively small following and low expectations, the MVP model can work wonderfully. You launch something as early as possible, you test your assumptions, you learn from your small but passionate audience, and then you iterate until you've got something extraordinary. Along the way, your (tiny) organization is associated with an ever-improving product, and by the time large groups of influencers and potential customers hear about you, you're in great shape to be perceived as a leader and innovator.

Conversely, if you already have a big following with high expectations, publicly launching a traditional MVP (one that leans more to the "minimum" side of the acronym than the "viable" side) can be disastrous. If you've reached a certain scale (which could vary depending on the reach of your organization versus the size of your field), perception and reputation are huge parts of your current and future success. A not-up-to-par product launch can hurt that reputation in the market and be perceived as a reason to avoid your company/product by potential customers. It can carry an MVP hangover for years, even if you do improve that product. And it can even drag down perception of your historic or currently existing products by association.

Tesla's a great example of an early-stage company that could not afford to launch an MVP. Prior to producing its first mass-market-available vehicle (the Tesla S in 2008), its reputation and reach was already so vast by virtue of Elon Musk's fame and the media surrounding the formation, growth, and, later, government loans to the company that anything less than extraordinary would have sundered its public perception and perhaps shuttered the organization.

Or take Slack, one of the darlings of the SaaS and tech world over the last few years. Had it become publicly associated with its initial MVP group chat product (which wasn't nearly as good, as feature-rich, or as compelling in user experience as HipChat or Yammer), it's likely that it never would have achieved the great success it did. But because Slack started small, with an internal-only product that slowly spread until it was truly exceptional and ready to earn broad adoption and reach, the model of iterating internally on a minimum viable product but waiting until that product was exceptional before launching worked wonders.

When Google launched in 1999, it could afford to be just like Slack—a slightly different and sometimes better version of what other search engines like Yahoo! and AltaVista were at the time. But today, if you wanted to compete with Google, your search engine would have to be immensely, obviously better in order to have even an iota of a shot against it. Microsoft's Bing, when it launched in 2009, sadly wasn't massively superior and thus, despite being an impressive effort, was generally perceived as a poor also-ran. Bing's gotten much better since then, and is now as good as or better than Google on most queries, but that MVP hangover has stuck with the brand for years and, in my opinion, continues to dampen the prospects of what should be a very decent option for web searchers.

The Alternative: An Exceptional Viable Product

My proposal is that we embrace the reality that MVPs are ideal for some circumstances but harmful in others, and that organizations of all sizes consider their market, their competition, and their reach before deciding what is *viable* to launch. I believe it's often the right choice to bias to the EVP, the "exceptional viable product," for your initial, public release.

It is absolutely the right move to first *build* an MVP. Developing every feature to perfection before you have anything people can test in the real world can be devastating. But as we saw in chapter 11, depending on your brand's size and reach, and on the customers and potential customers you'll influence with a launch, I'd urge you to consider whether a private launch of that MVP, with lots of testing,

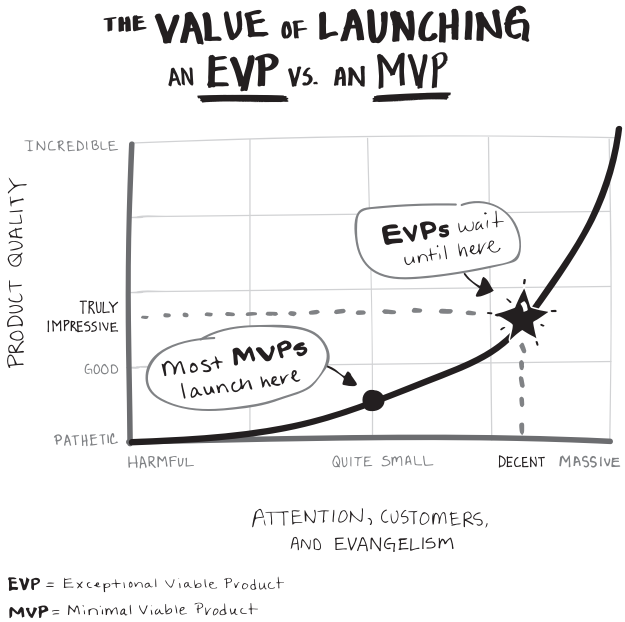

learning, and iteration to a smaller audience that knows they're beta testing, could be the best path. For Moz, it's worked out remarkably well in several of our product efforts. I'll share an example that highlights how that process can work.

After the disastrous launch of Moz Analytics and the mediocre reception of Spam Score, I vowed not to make the MVP mistake again. In early 2015, I pitched Moz's executive team on a plan to build a new keyword research tool. Despite my last few product challenges, I got approval to work with a small team of five engineers on my proposal, with the caveat that those folks would have to continue their regular work maintaining the infrastructure, upgrades, and operations of another product. I talked to the team, and they agreed they could simultaneously support our existing workload and take on this new project.

I knew that if I messed this up, it would likely be my last shot at having product ownership responsibilities at Moz. I can even recall the exact words my CEO had with me as I was embarking on it. I asked if she trusted me to do a good job on these new keyword research tools, and she replied: "I think you're good at seeing what's wrong with a product, but I don't think you're a good product designer." Those words became my mantra and my motivation over the next year. I was determined to prove to her, to my team, and to myself that I'd learned my lessons and could still make something remarkable.

And yes, I realize that proving other people wrong is not an emotionally healthy form of motivation, but for me, it's always been an effective one.

The new product I wanted to build would eventually be called "Keyword Explorer." The goal was to help people determine what their customers and audiences were searching for online, craft lists of keyword terms and phrases to target on their websites, and prioritize those lists so they could logically know what to work on first. For example, if you're starting up a new Italian foods website, and you want to help your visitors choose the best brands of fresh and dried pasta, you need to know what people type into Google when they're looking for that type of content. You also want to know which words and phrases are used more commonly than others—do more people search for "fresh pasta" or "dried pasta"? More for "best pasta brands" or "best pasta makers"? And which of those search queries are harder versus easier to appear in (usually determined by which other sites are already ranking in the results and how popular/important/well-linked-to they are).

My first step was to create a slide deck in which I "pitched" the concept of a keyword research tool to the Moz team. I shared this first with my executive team and then with the entire company. I made three primary claims about why we should invest in solving the problem of keyword research:

1. The keyword research tools space was far less competitive at the time than other areas in which Moz had invested product and engineering resources (like links, rankings, and website crawling).
2. Keyword research was (and still is) the most commonly performed task in the SEO world, mostly because it's important to do every time you produce new content.
3. I believed that Moz had access to unique data and the ability to combine that data in ways that no other company in the world could.

My research process for the pitch and the product design involved a lot of requests to my friends and colleagues in the SEO community. I asked a few dozen folks for detailed information about how they did keyword research currently and what tools and processes they applied. I sent out a pair of surveys via Twitter and BCC email blasts to get a sense of what was already popular in the field (and wrote a blog post about it). Those surveys showed me which tools people were most familiar with and which ones people paid for. That gave us a huge head start in knowing what naming conventions and design layouts and existing structures our audience was already familiar with.

And, finally, I sat down, face-to-face, with a handful of SEO practitioners as they showed me exactly how they did keyword research for their websites or clients.

Those surveys, interviews, and email conversations confirmed a few big truths:

- Almost everyone who did SEO professionally used a very similar methodology for keyword research. First, they gathered a bunch of potential terms and phrases from lots of different online tools. Next, they put them all in a spreadsheet (Excel or Google Sheets) and, often manually, gathered metrics about the search popularity, relative ranking difficulty, and traffic potential. Finally, they used some version of a formula to combine the metrics and prioritize the list.

- No existing tool performed this work automatically. Everyone had to use multiple tools and data sources alongside a manual process to build their lists.
- A strong majority of professionals were paying for one or more keyword research tools.
- The accuracy (or rather, gross inaccuracy) of keyword volume data was perceived as the biggest challenge and frustration among people I surveyed and talked to.

My one fear coming out of this process was that I hadn't learned much, if anything new. All the data I'd collected and all the interviews I did merely confirmed my preexisting suspicions. I worried that I'd somehow biased my product research—it couldn't be that my intuition was perfectly matched with reality, right?

Well, maybe. In this case, I wasn't just a product designer, I was also technically my own customer. I create lots of content for the web, and I'd been doing keyword research myself for more than a decade. Like a lot of the folks I spoke to and surveyed, I, too, used many tools, created spreadsheets, and built my own prioritization calculations. I went to conferences and events, read blog posts, and watched industry insiders and newcomers write about, talk about, tweet about, and share their processes. I lived and breathed the life of my customers, and, I think because of that, my empathy radar was especially well tuned to this problem set. I'm still very, very glad I did the research, though, because I needed it to prove to my team what was necessary for a release-worthy tool.

We started building Keyword Explorer in May 2015, initially hoping to launch by the end of the year. That slipped into January, at which point we had a workable version that included the basics. Input a word or phrase and it would analyze Google's search results, suggest related search queries, return volume estimates, and enable you to easily build a list for export or analysis. I was pretty excited about the tool and went about showing it off to many of my SEO colleagues. I stayed over at Wil Reynolds's house (my CEO swap friend) in January and walked him through the tool. I shared it with colleagues from Distilled, another consulting company to whom Moz has had long-time close ties. But I think the experience that triggered my change of heart came most saliently from a demo I gave to Dan Shure.

Dan's an SEO consultant and well-regarded podcast host (of *Experts on the Wire*) based in Worcester, Massachusetts. He has a sizable following in the SEO community and is particularly passionate about keyword research. Dan and I were scheduled to be at a conference together in Orlando, Florida, at the end of January, just before Keyword Explorer was scheduled to launch.

When I opened my laptop and showed Dan our application, I was nervous. I desperately wanted him to like what we'd built, and I knew that if he didn't, we'd have serious work ahead. I typed in a few keywords, showed him the basics, and invited him to try some of his own searches.

In the good news column, the interface was intuitive. Dan picked it up right away and started building a list of keywords for one of his recent clients. He had a lot of questions about our data sources and whether he could trust the numbers we provided. He seemed mostly satisfied by my explanations but said he'd need to double-check them against his own data and the Google search advertising campaigns his client was running.

In the bad news column, Dan was underwhelmed by a number of missing features, a lack of clarity around the sources of some of the data, and the inability to filter, sort by, and view metrics he deemed critical to his keyword selection process. I showed him some of the behind-the-scenes data sources (essentially command-line-style interfaces that could surface some of the filters and information he wanted), and his eyes got wide with excitement.

That night, I wrote up a lengthy email sharing the notes from my review session with Dan. I told the development team, and in a later email my executive team, that based on Dan's feedback (and several others'), I wanted to delay the launch until we had released five more features I felt were critical to earning the reception and adoption we wanted. Frustratingly, those features would take an additional four months of development! We were going to be nearly six months late delivering the product launch.

Thankfully, because Moz had so many other bets placed for 2016, and because our small team and the Keyword Explorer tool was considered a side project less central to the company's growth plan for the year, we were granted the reprieve. From January to May, we worked tirelessly to upgrade and improve the tool. When we finally launched, it was with a product of which we could be truly proud.

Those first two days after launch, we had just over seventy thousand visits to Keyword Explorer from more than sixteen thousand unique visitors. Keyword Explorer was the second-most up-voted product on Product Hunt the day it launched and made news on more than a dozen industry publications in the SEO and web-marketing fields. Hundreds of people commented on the launch post and on social media sites, nearly all of them in an overwhelmingly positive fashion.

My favorite came from Dan Shure himself, whose endorsement I wasn't sure we'd earn given his critiques back in January:

> *Incredibly excited to see @Moz's new Keyword Tool launched. It's a "must-use" for all my keyword/topic research now.*
>
> **—Dan Shure's tweet on the day of Keyword Explorer's public launch**

Over the next year, Keyword Explorer would become the most lauded, positively reviewed, and fastest-growing new feature (by use) in Moz's subscription. Our customer success team shifted much of their onboarding efforts to showcase the product, and our metrics showed that its use by customers had a strong correlation with subscribers who stuck around as paying customers versus those who quit.

It wasn't an unmitigated success, though. We offered access to Keyword Explorer two ways: through the traditional month-to-month Moz subscription (at $149/month) or as a stand-alone product (which could only be purchased with an up-front payment of $600 or a higher-tier at $1,800). The month-to-month subscriptions were significantly bolstered by the launch and the product's success, but the stand-alone subscriptions had almost no adoption at all. Don't underestimate the power of packaging and pricing.

For me, the lesson about MVPs versus EVPs was a powerful one. In the future, I don't think I'd ever willingly publicly release a product that leans minimum rather than exceptional (at least, not with a brand that has any substantial reach). The impact of reputation and word of mouth on potential success versus the risks of MVP hangover is too important to ignore.

9780735213326_Lost_TX.indd 192 7/19/18 11:16 AM